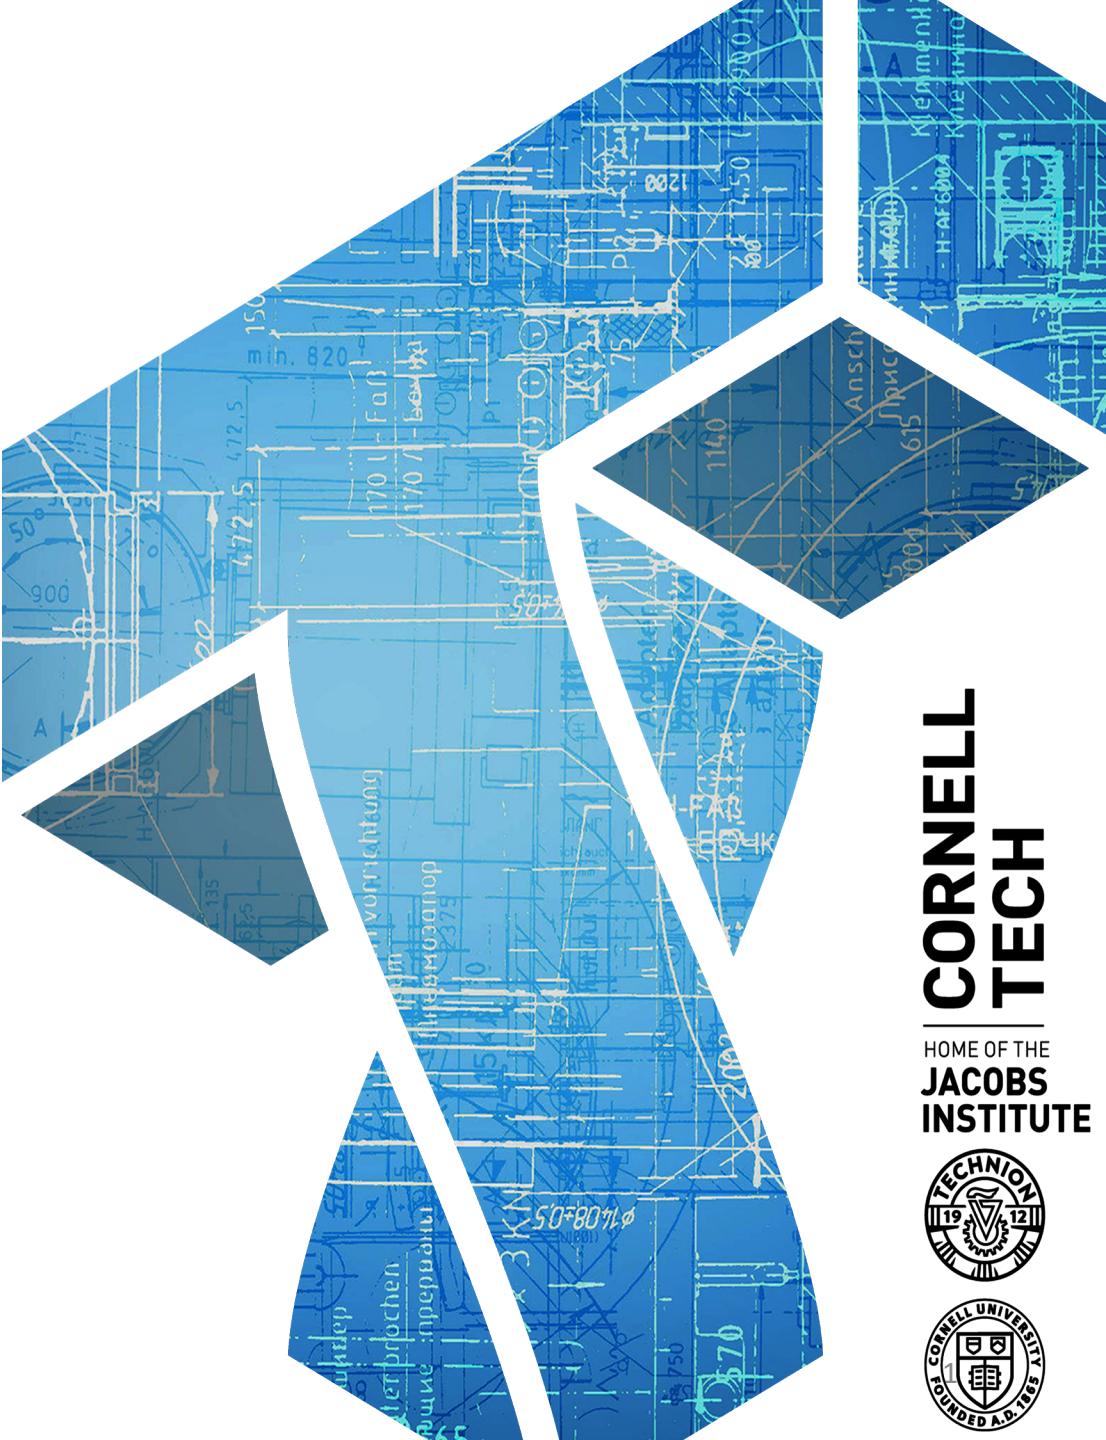


# CS 5435: OS security

Instructor: Tom Ristenpart

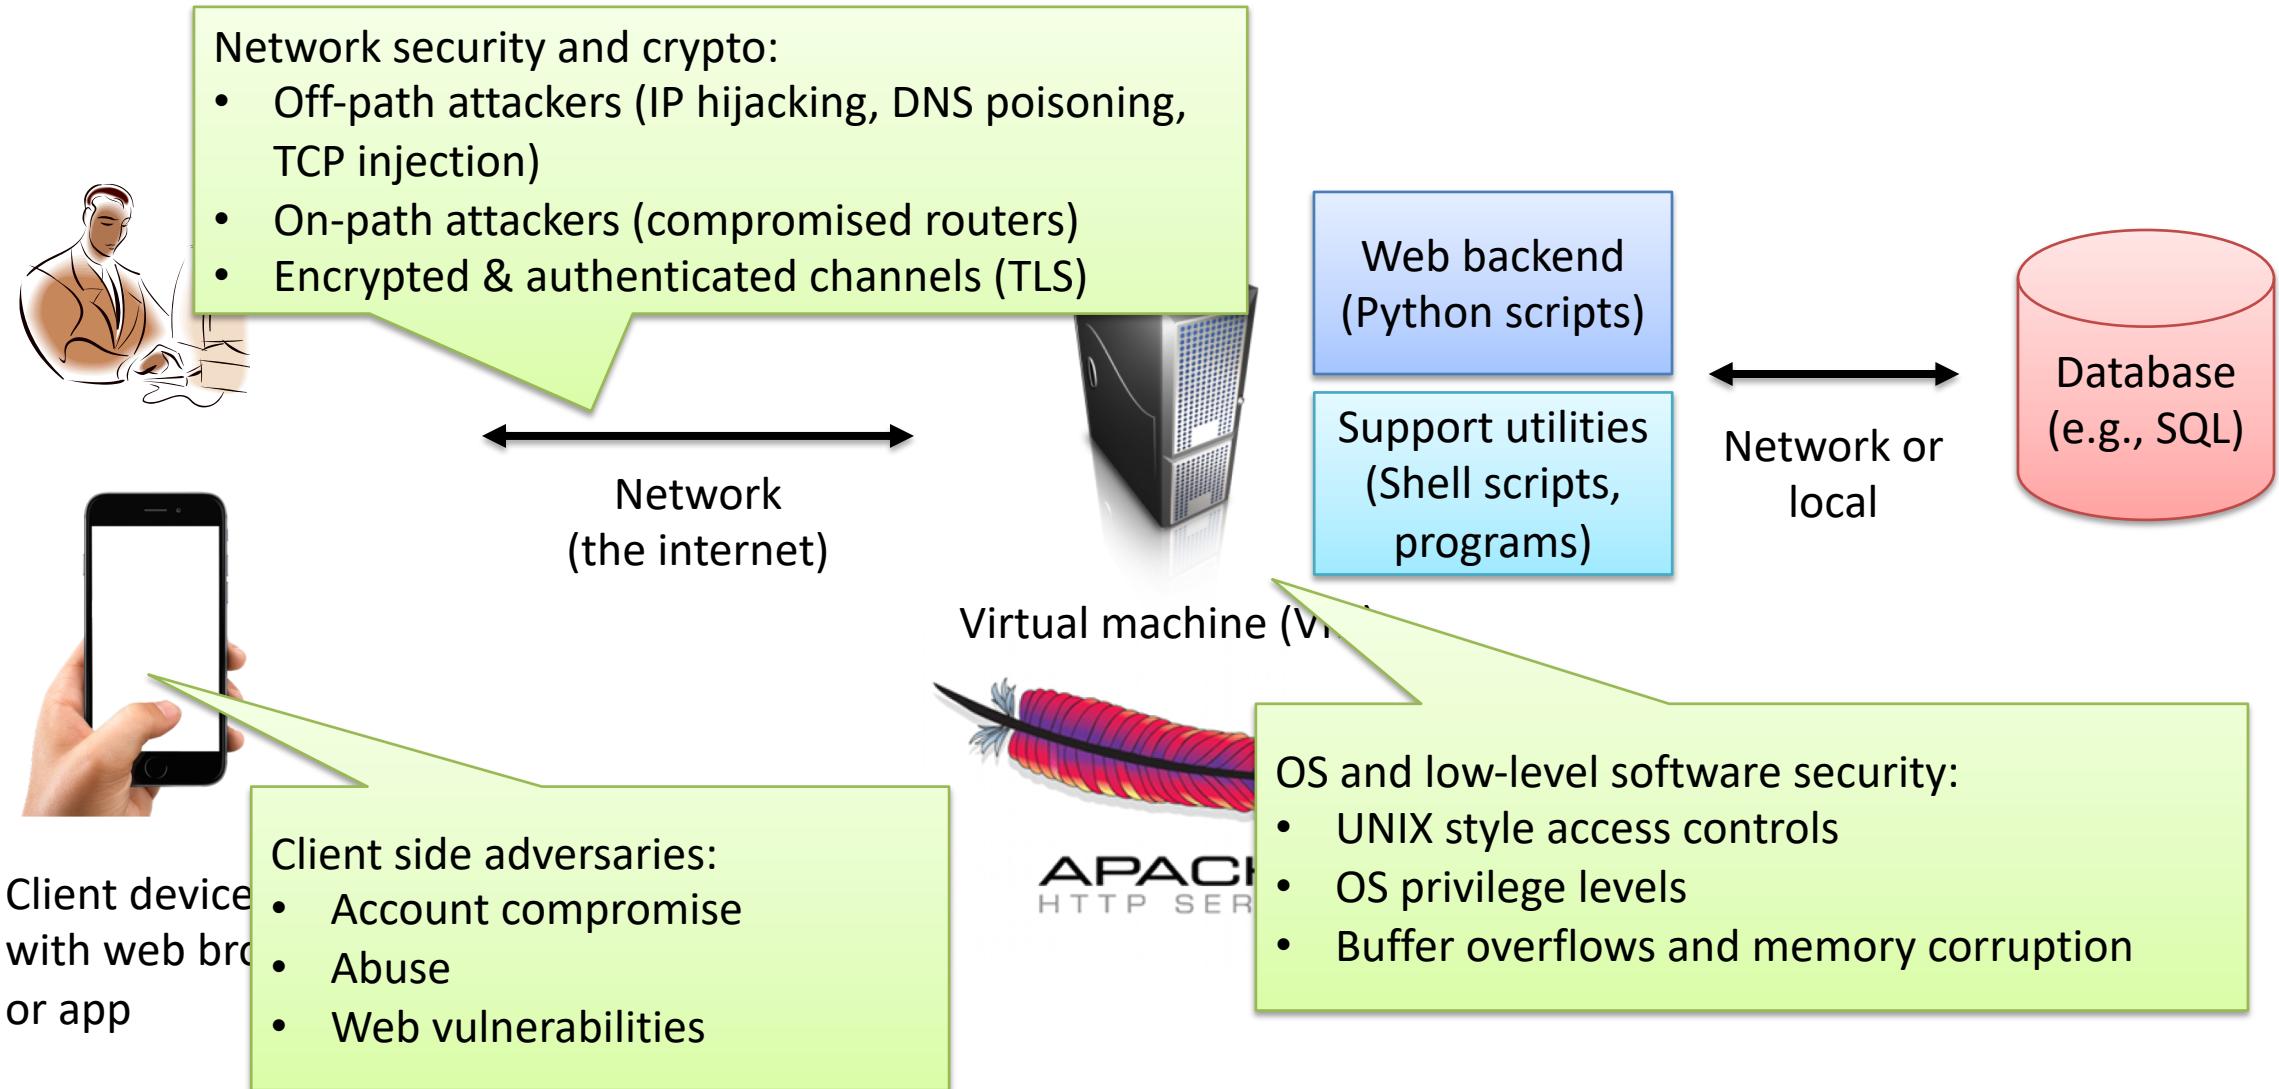
<https://github.com/tomrist/cs5435-fall2019>



# Announcements

- HW3 due date extended to Monday, Nov 11
- Office hours next couple days
- HW4 will be released later this week
- No lecture on Monday Nov 11 (holiday)
- Possibly no lecture on Wednesday Nov 13, may do review session for HW4

# Where we're at via our web service example



# Take yourself back to the 1960's...

Time-share multi-user computers  
coming into use

GE-645  
36 bit address space  
Up to 4 processors  
Magnetic tape drives  
Supported virtual memory in hardware



Courtesy of  
<https://www.bell-labs.com/var/articles/invention-unix/>

# Multiplexed Information and Computing Service (Multics)

Project to develop operating system for time-shared systems

- Designed from 1964-1967.
- MIT project MAC, Bell Labs, and GE
- ~100 installations at greatest extent
- Last one shut down in 2000 (Canadian department of defense)

“A small but useful hardware complement would be 2 CPU units, 128K of core, 4 million words of high speed drum, 16 million words of disc, 8 tapes, 2 card readers, 2 line printers, 1 card punch and 30 consoles.”

[Vyssotsky, Corbato, Graham 1965]

# Multics: ancestor to many OS's

Lots of innovations in design

- Use of segmentation and virtual memory with hardware support
- SMP (shared memory multiprocessor)
- Written in PL/1 (high level language)
- Inspired developers of Unix

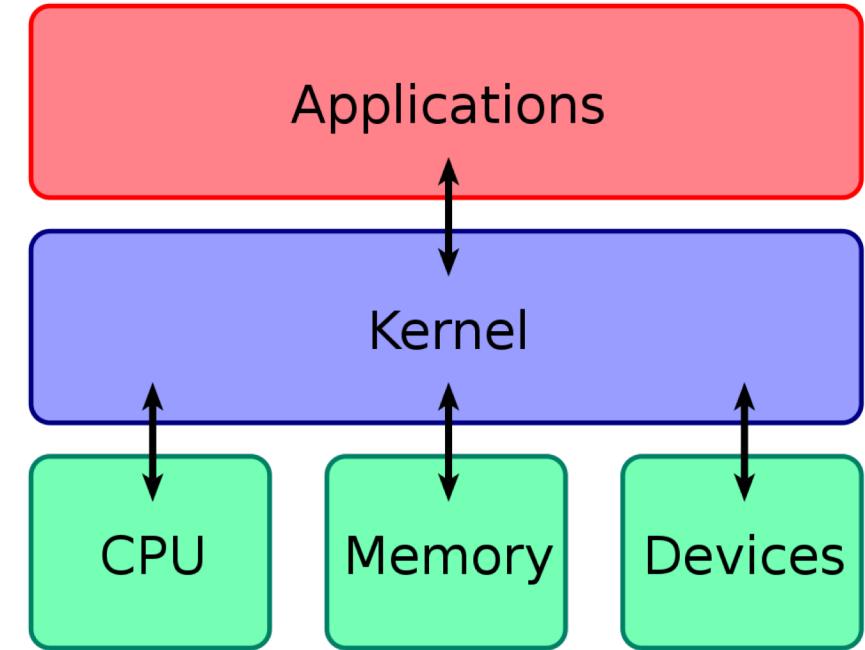


F. Corbato

Significant attention paid to security

# Recall operating system basics

- Multi-tasking, multi-user OS are now the norm
- ***Kernel*** mediates between applications and resources
- ***Applications*** consist of one or more processes
- ***Processes*** have executable program, allocated memory, resource descriptors (e.g., file descriptors), processor state

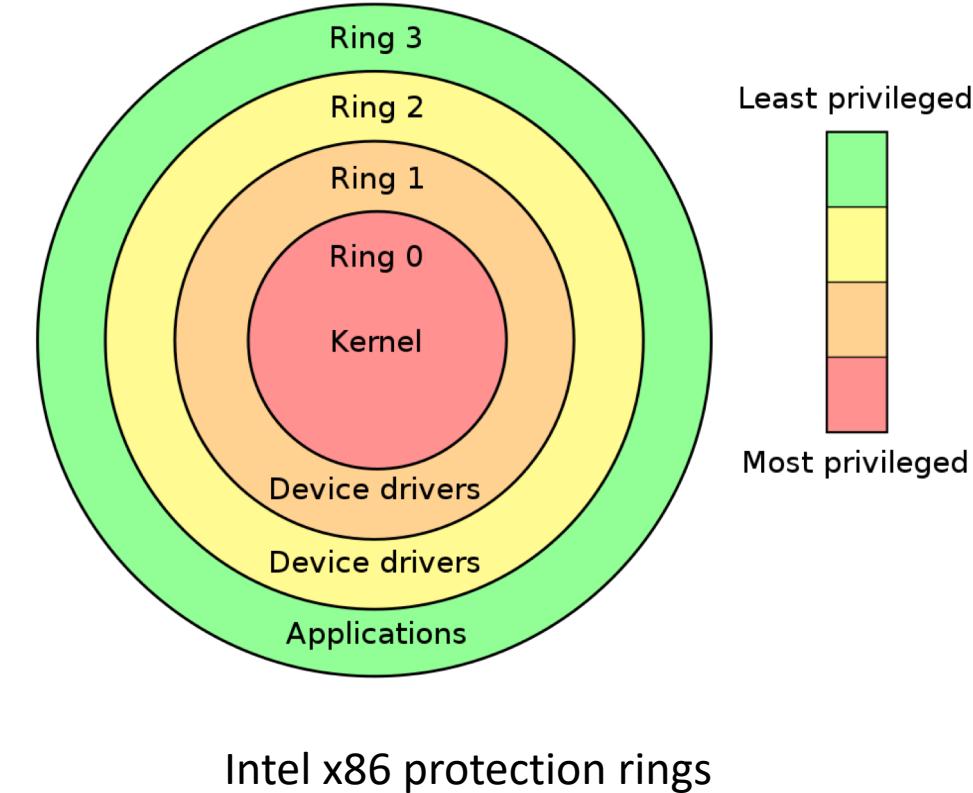


# Privilege levels and protection rings

Different parts of system must operate at different privilege levels

Protection rings included in all typical CPUs today and used by most operating systems

- Lower number = higher privilege
- Ring 0 is supervisor
- Inherit privileges over higher levels



Intel x86 protection rings

## ***Principle of least privilege:***

User account or process should have least privilege level required to perform their intended functions

# Systems security ingredients

- Security model:
  - Abstraction of system to help us express security policies
- Security policies:
  - Specification of what parts of system should be allowed to do
- Security mechanism:
  - How we implement the policy

# Systems security ingredients

- Security model:
  - Subjects
  - Objects
  - Operations

	Unix	Web
Subjects (Who)	Users, Processes	Domains
Objects (What)	Memory, Files, Hardware devices ...	DOM components, cookies, ...
Operations	Read, write, execute	Read, write, ...

# Access control matrix

Most frequent way of specifying security policy

		Objects		
		file 1	file 2	...
Subjects	user 1	read, write	read, write, own	read
	user 2			
	...			
	user m	append	read, execute	read,write, own

User i has permissions for file j as indicated in cell [i,j]

Due originally to Lampson in 1971

# Two common implementation paradigms

	file 1	file 2	...	file n
user 1	read, write	read, write, own		read
user 2				
...				
user m	append	read, execute		read,write, own

## (1) Access control lists

Column stored with file

## (2) Capabilities

Row stored for each user

Unforgeable tickets given  
to user

Discretionary access controls: users can set some access controls (e.g., Unix file system ACLs)

Mandatory access controls: access controls set in one centralized location (e.g., SELinux)

# Unix Security Model

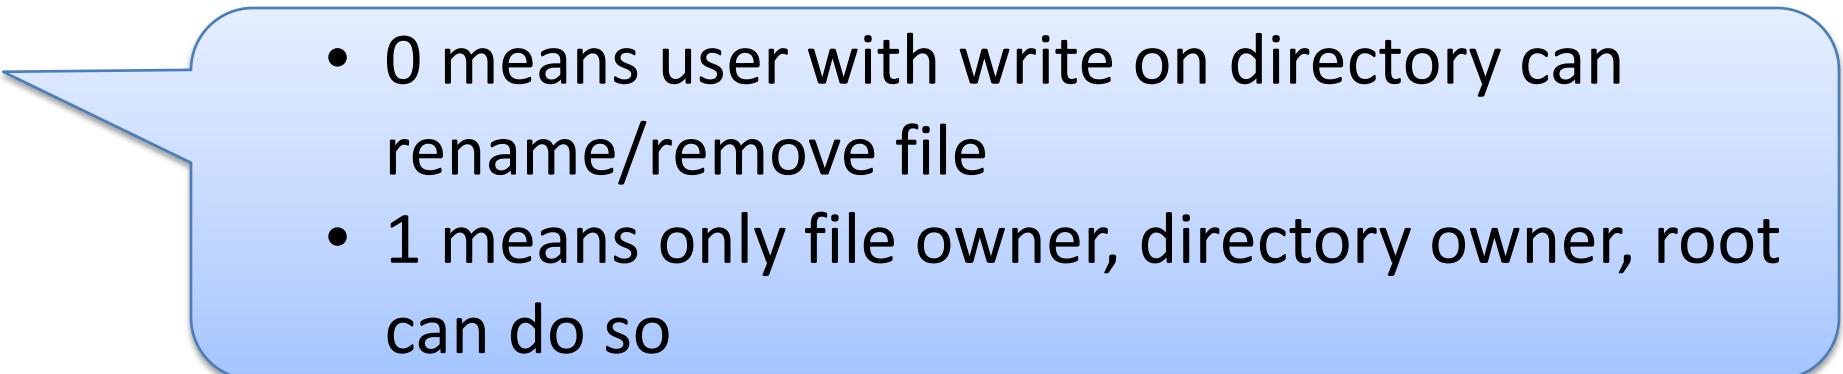
- Users include high-privilege *root* (*aka superuser*)
  - User ID (UID) is 0
- Users authenticated traditionally via passwords
  - /etc/passwd
  - /etc/shadow
    - Contains password hashes, restricted access

	Unix
Subjects (Who)	Users, Processes
Objects (What)	Memory, Files, Hardware devices ...
Operations	Read, write, execute

# Unix Security Model

- File system permissions uses discretionary ACL approach
- Each file has 12 ACL bits
  - rwx for each of owner, group, all
  - Set user ID
  - Set group ID
  - Sticky bit

	Unix
Subjects (Who)	Users, Processes
Objects (What)	Memory, Files, Hardware devices ...
Operations	Read, write, execute

- 
- 0 means user with write on directory can rename/remove file
  - 1 means only file owner, directory owner, root can do so

# UNIX-style file system ACLs

rist@seclab-laptop1.local: ~/work/revindiff/full — less — 80x24

```
total 27648
drwxr-xr-x  51 rist  staff      1734 Aug 23 13:11 .
drwxr-xr-x  46 rist  staff      1564 Jul  5 12:37 ..
drwxr-xr-x   7 rist  staff      238 Jun 22 18:29 .svn
-rw-r--r--   1 rist  staff      321 Jun  2 22:38 Makefile
-rwxr-xr-x   1 rist  staff    258319 May 11 00:18 abbrev.bib
-rwxr-xr-x   1 rist  staff    242609 May 11 00:18 abbrev_short.bib
-rw-r--r--   1 rist  staff    3049 Jun 20 14:22 abstract.tex
-rw-r--r--   1 rist  staff    6921 May 11 00:18 accents.sty
-rw-r--r--   1 rist  staff     534 Jun 20 16:30 acknowledgements.tex
-rw-r--r--   1 rist  staff      100 Jun 20 16:30 acknowledgements.tex.bak
-rw-r--r--   1 rist  staff      100 Jun 20 16:30 acknowledgements.tex.bak
Permissions:
- Directory?
- Owner (r,w,x) , group (r,w,x), all (r, w, x)
Owner (rist)
Group (staff)
```

The terminal window shows a file listing with a callout box highlighting the 'Owner' and 'Group' fields. A red oval surrounds the 'Owner (rist)' and 'Group (staff)' text in the callout box.

# Roles (groups)

Group is a set of users

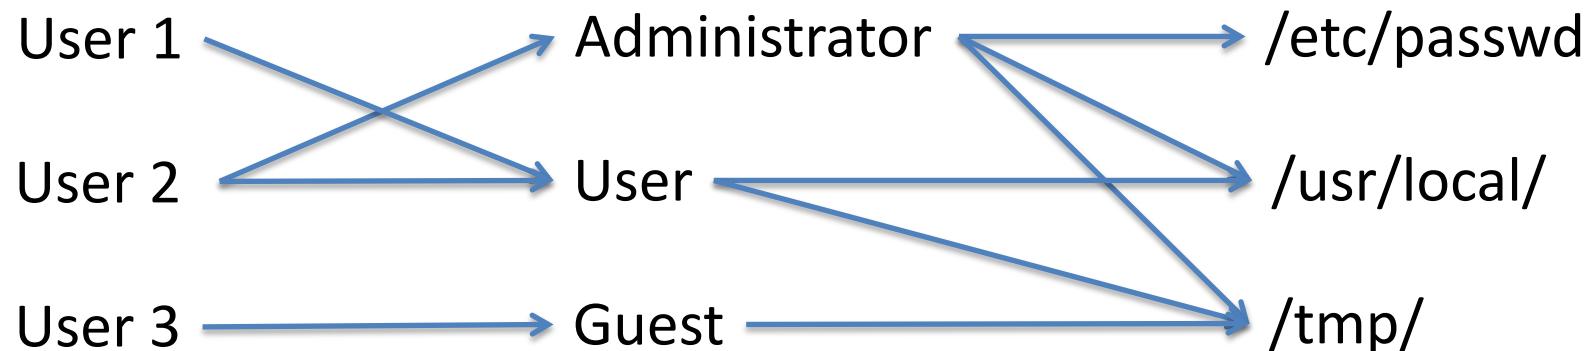
Administrator

User

Guest

Simplifies assignment of permissions at scale

Concept sometimes referred to as role-based access control (RBAC)

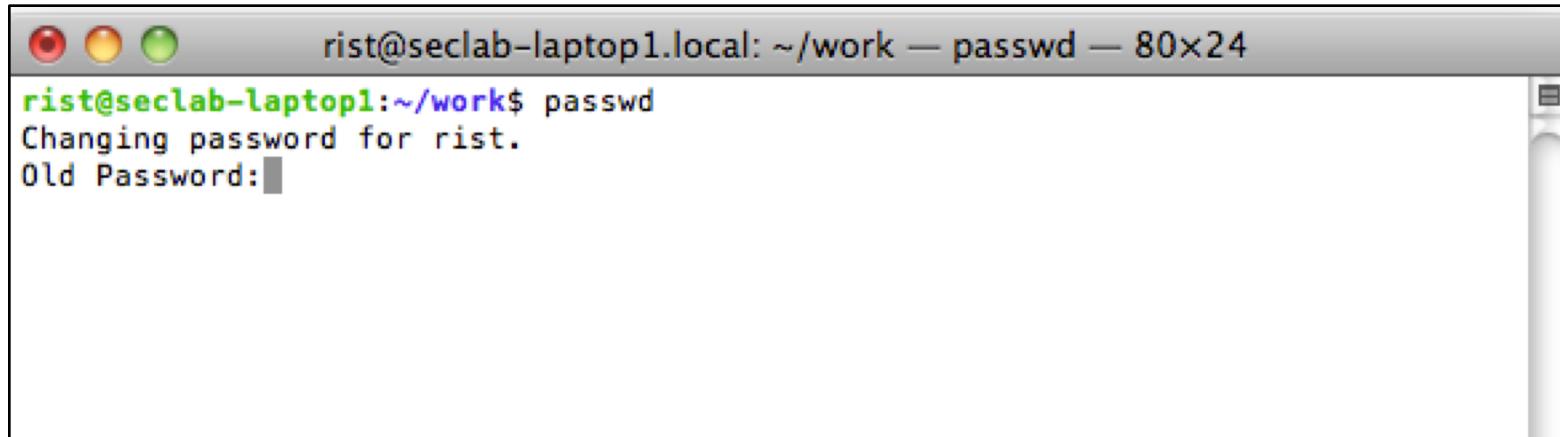


# UNIX file permissions

- Owner, group
- Permissions set by owner or root
- Resolving permissions:
  - If user=owner, then owner privileges
  - If user in group, then group privileges
  - Otherwise, all privileges

# UNIX Process permissions

- Process (normally) runs with permissions of user that invoked process



A screenshot of a terminal window titled "rist@seclab-laptop1.local: ~/work — passwd — 80x24". The window shows the command "passwd" being run by the user "rist". The terminal output includes "Changing password for rist." and "Old Password: [redacted]". The terminal has a standard OS X-style title bar with red, yellow, and green buttons.

/etc/shadow is owned by root

Users shouldn't be able to write to it generally

How can passwd reset user's password?

# Setuid programs

```
-rwxr-xr-x    1 root    wheel      4954 Feb 10  2011 znew
-r-xr-xr-x    1 root    wheel     63424 Apr 29 17:30 zprint
rist@seclab-laptop1:/usr/bin$ ls -al passwd
-r-sr-xr-x  1 root    wheel  111968 Apr 29 17:30 passwd
rist@seclab-laptop1:/usr/bin$
```

- setuid bit – execute with privileges of file's owner
- setgid bit – execute with privileges of file's group
- So passwd is a ***setuid program*** program runs at permission level of owner, not user that runs it

Least privilege at *process granularity*:  
passwd runs as root to access /etc/shadow  
Can we do better?

# Process permissions

Process (normally) runs with permissions of user that invoked process  
Unless executable is setuid or setguid

## Real user ID (RUID) --

same as UID of parent (who started process)

Used to figure out who started process

passwd example:  
RUID = rist

Also SGID, EGID, etc..

## Effective user ID (EUID) --

from set user ID bit of file being executed or due to sys call

indicates current permissions of process

passwd example:  
EUID = root

## Saved user ID (SUID) --

place to save the previous UID if one temporarily changes it

# seteuid system call

```
uid = getuid();
eid = geteuid();
seteuid(uid);    // Drop privileges
...
seteuid(eid);    // Raise privileges
file = fopen( "/etc/shadow", "w" );
...
seteuid(uid);    // drop privileges
```

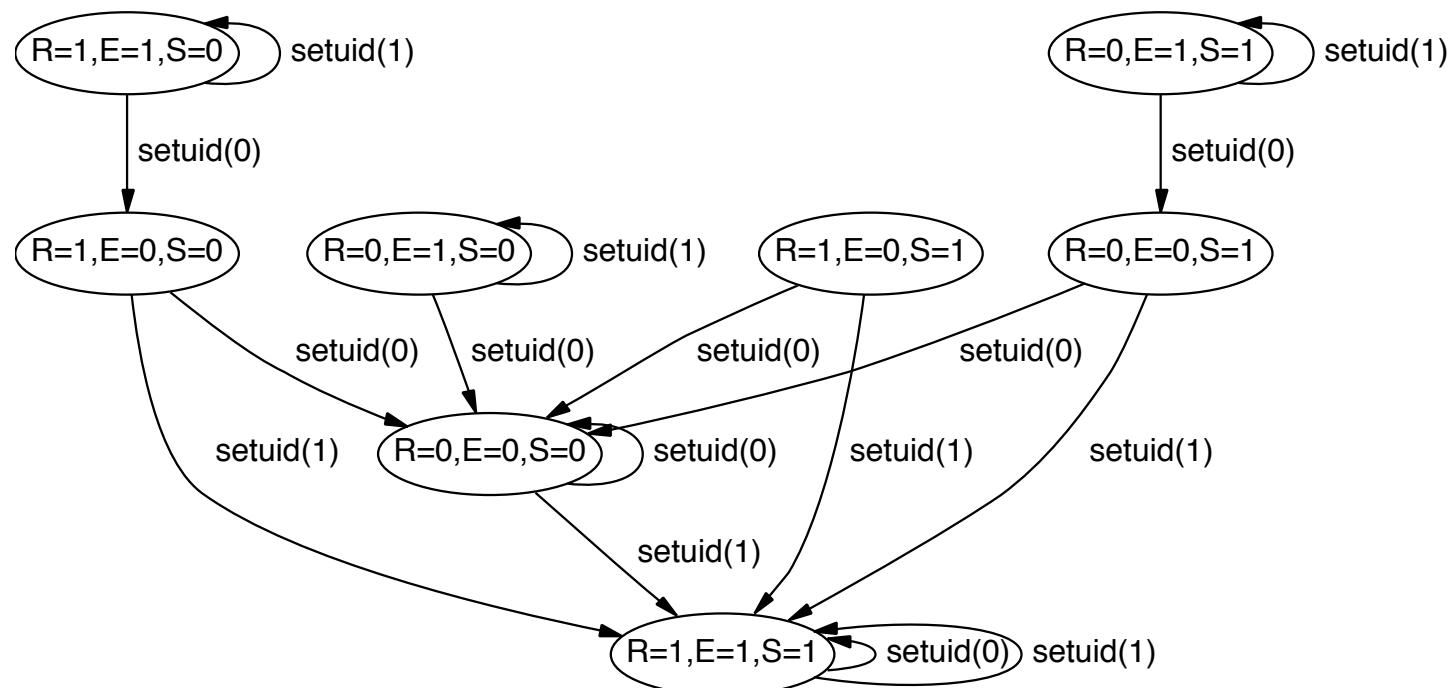
seteuid can:

- go to SUID or RUID always
- any ID if EUID is 0

Note: this is C pseudocode  
C still used widely for OS's and  
low-level programming

# Details of setuid more complicated

Chen, Wagner, Dean “Setuid Demystified”



(a) An FSA describing *setuid* in Linux 2.4.18

# Setuid allows necessarily privilege escalation but...

- Source of many ***privilege escalation vulnerabilities***
  - Bug that allows lower-privilege user to perform actions as higher-privilege user (most often: root)

Race conditions

Control-flow hijacking vulnerabilities in local setuid program gives privilege escalation

# Race conditions

## Time-of-check-to-time-of-use (TOCTTOU)

```
if( access(“/tmp/myfile”, R_OK) != 0 ) {  
    exit(-1);  
}  
  
file = open( “/tmp/myfile”, “r” );  
read( file, buf, 100 );  
close( file );  
print( “%s\\n”, buf );
```

access checks RUID,  
but open only checks EUID

Say program is setuid root:  
access checks RUID, but open only checks EUID

Time

access("/tmp/myfile", R\_OK)



In -s /etc/shadow /tmp/myfile

open( "/tmp/myfile", "r" );

print( "%s\n", buf );

Prints out shadow file  
(including password hashes)

# Better code

```
euid = geteuid();
ruid = getuid();
seteuid(ruid);          // drop privileges
file = open( "/tmp/myfile", "r" );
read( file, buf, 100 );
close( file );
print( "%s\n", buf );
```

# Low-level software vulnerabilities

```
int foo() {  
    int a = 0;  
    return a + 7;  
}
```

C code

Compiler

```
pushl %ebp  
movl %esp, %ebp  
subl $16, %esp  
movl $0, -4(%ebp)  
movl -4(%ebp), %eax  
addl $7, %eax  
leave  
ret
```

Disassembled machine code

Control flow refers to sequence of instructions followed by CPU

Control flow hijacking:

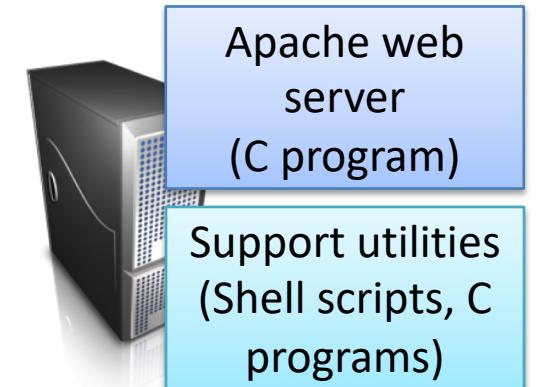
- exploiting vulnerability to adversarially manipulate control flow

# Low-level software vulnerabilities

C programs notorious for vulnerabilities that allow *control flow hijacking*

## Local privilege escalation:

- Given user-level access, obtain root-level access
- Typical exploit gives attacker root shell



## Remote exploit:

- Gain illicit access remotely over the network
- Typical exploit gives attacker remote shell

# Summary

- Multics: seminal multi-user operating system
  - many security features
  - significant auditing performed, achieved high security certifications
- Security models, policies, mechanisms
  - Subjects, Objects, Operations
  - Access control lists
- Unix security
  - File system permissions
  - Setuid programs
- Race conditions, intro to control-flow hijacking

