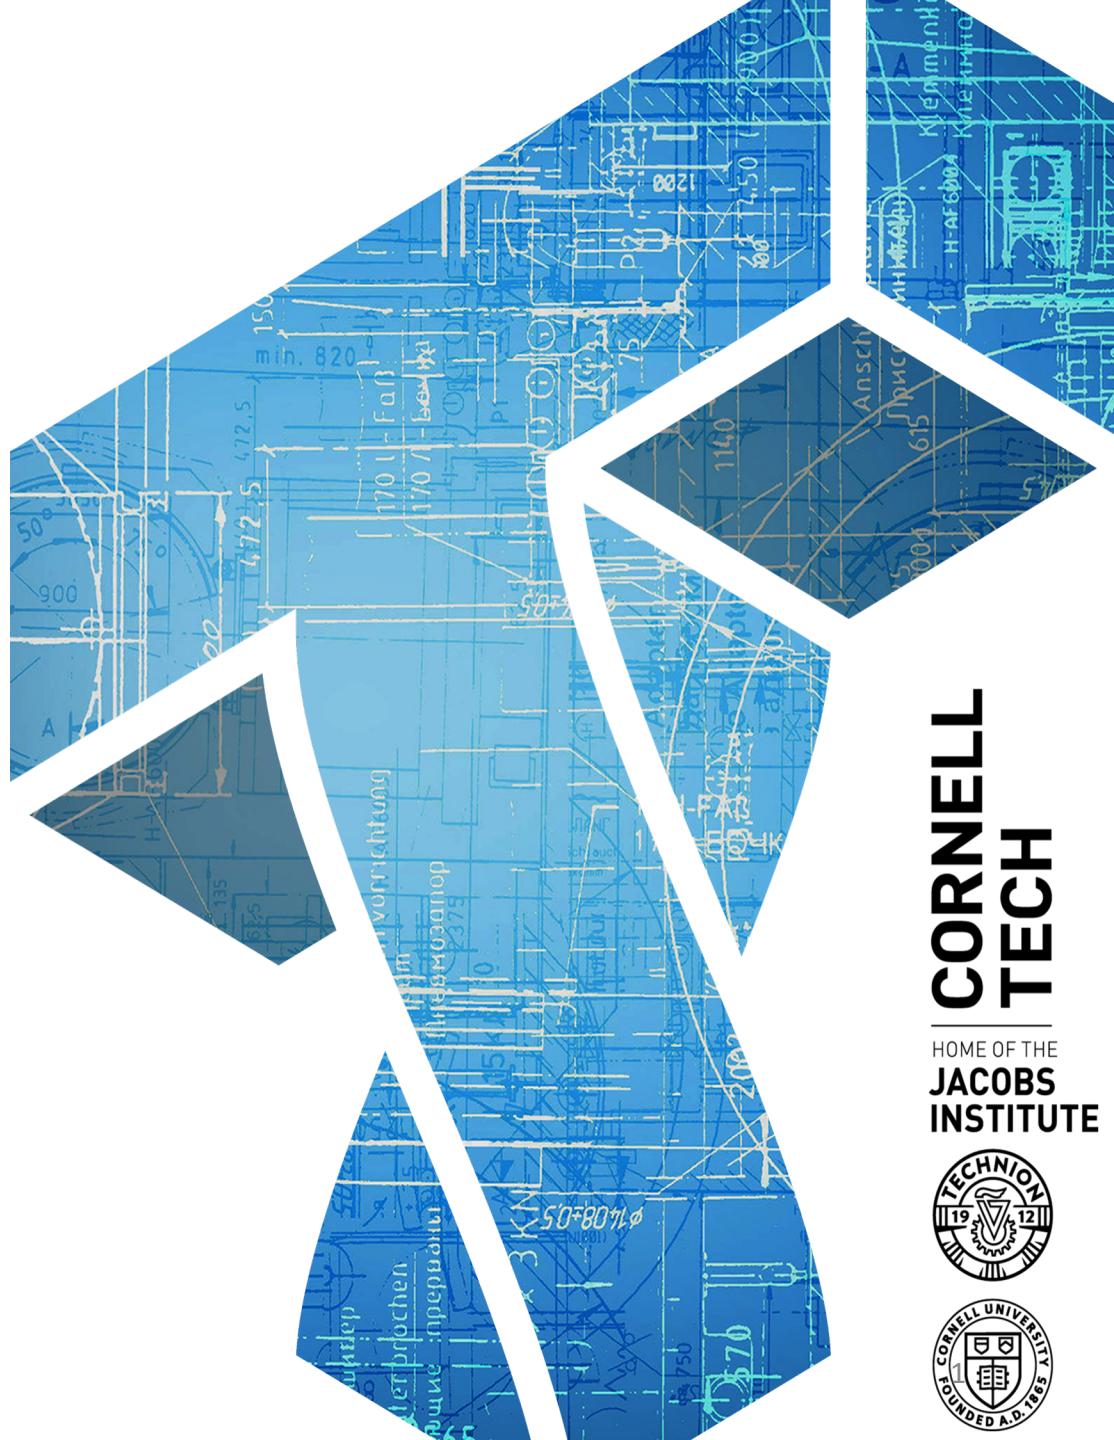


# CS 5435: Memory protection mechanisms

Instructor: Tom Ristenpart

<https://github.com/tomrist/cs5435-fall2019>



**CORNELL  
TECH**

HOME OF THE  
**JACOBS  
INSTITUTE**



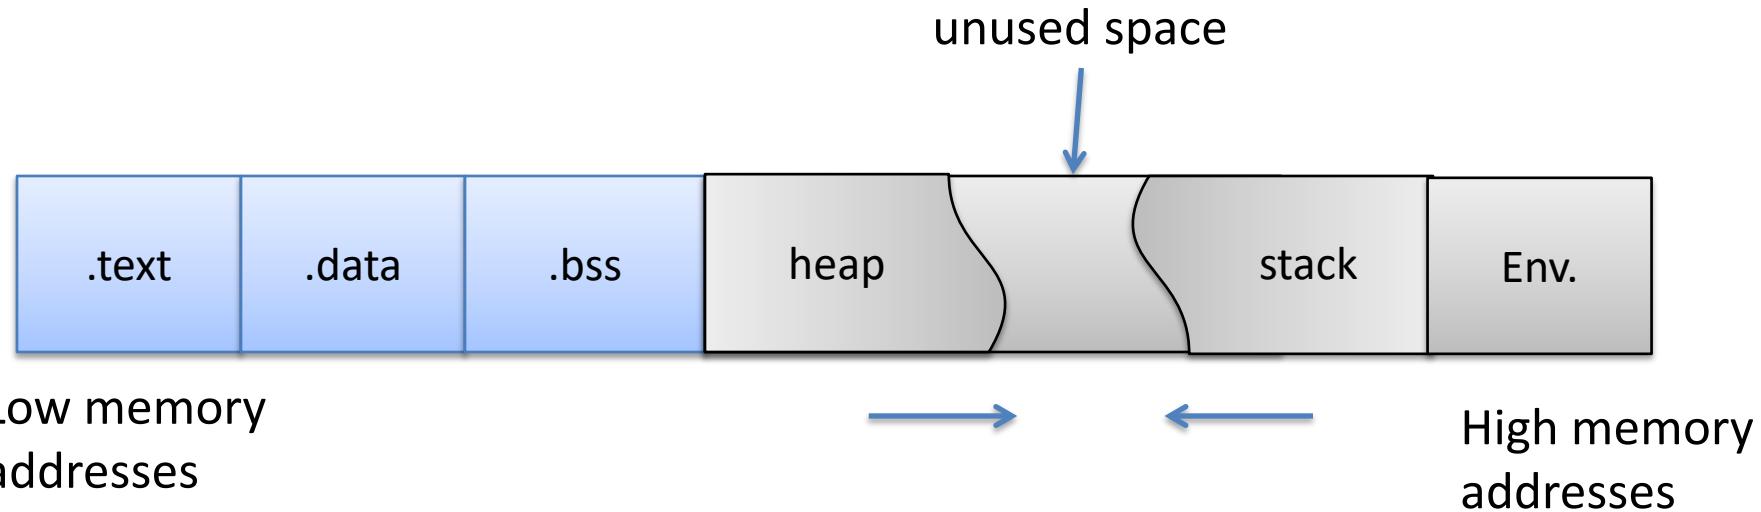
# Running demo example

(modified from Gray hat hacking book linked in resources)

```
1 #include <stdio.h>
2 #include <string.h>
3
4
5 void greeting( char* temp1 )
6 {
7     char name[400];
8     memset(name, 0, 400);
9     strcpy(name, temp1);
10    printf( "Hi %s\n", name );
11 }
12
13
14 int main(int argc, char* argv[])
15 {
16     greeting(argv[1]);
17     printf( "Bye %s\n", argv[1] );
18 }
```



# Process memory layout



.text:  
machine code of executable

.data:  
global initialized variables

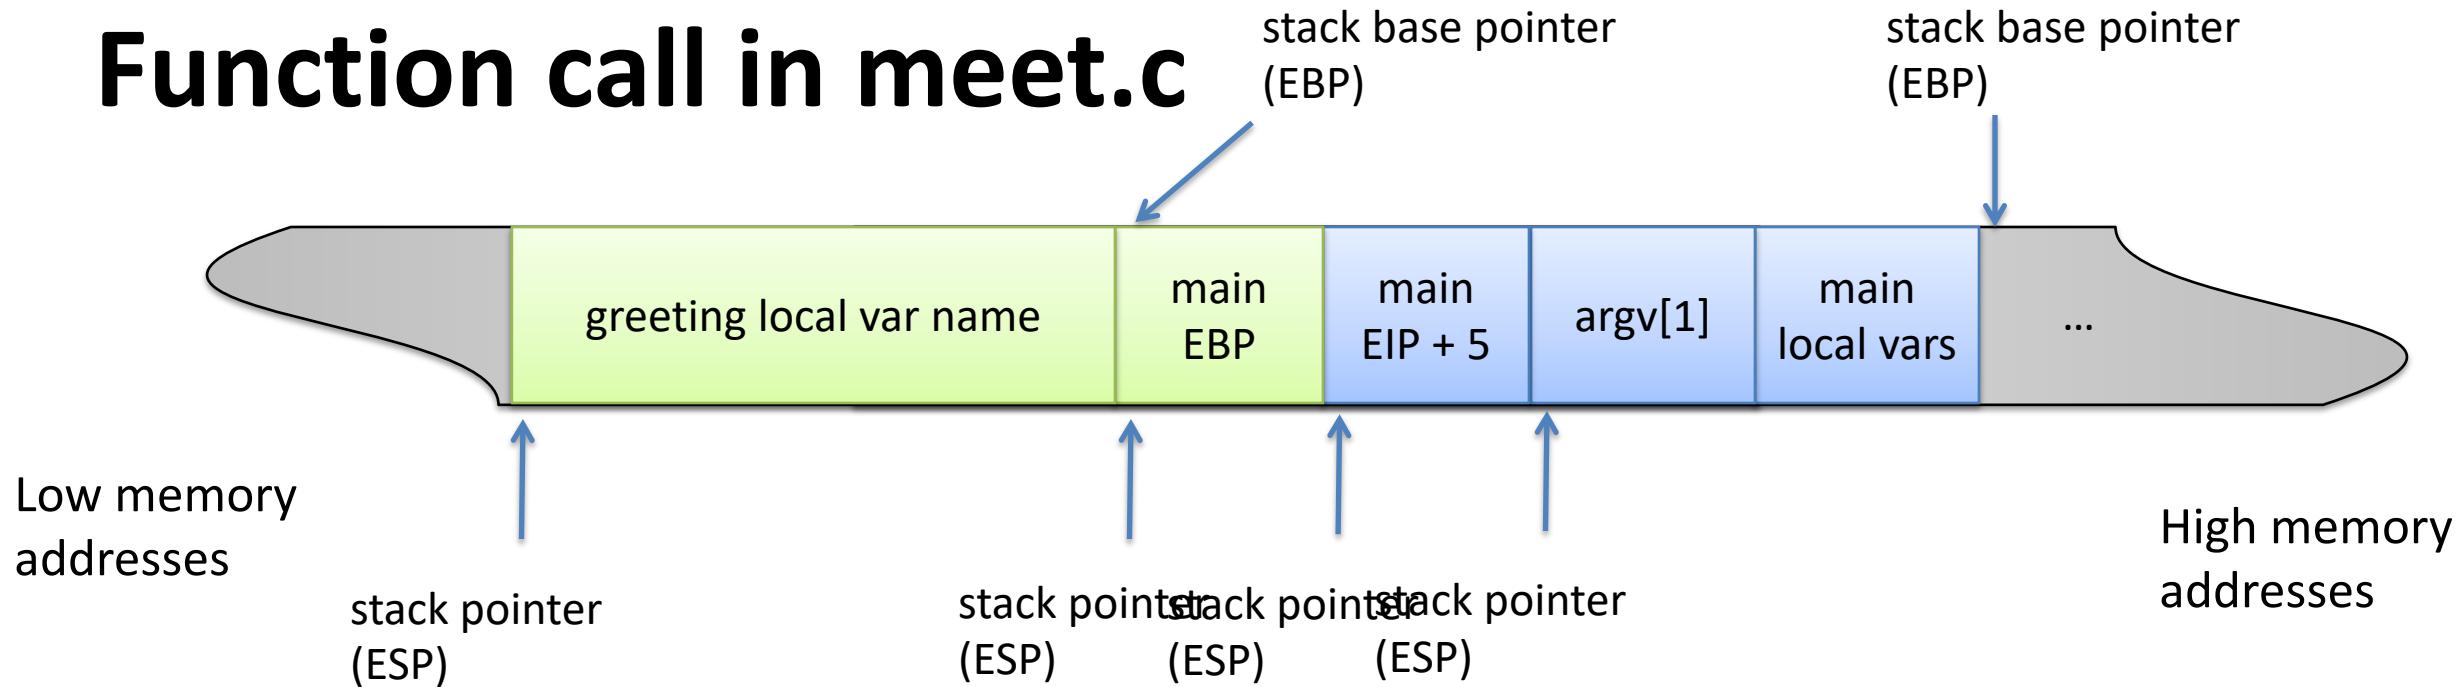
.bss:  
“below stack section”  
global uninitialized variables

heap:  
dynamic variables

stack:  
local variables, track func calls

Env:  
environment variables,  
arguments to program

# Function call in meet.c



```
student@5435-hw4-vm:~/demo$ gdb -q meet
Reading symbols from meet...done.
(gdb) disassemble main
Dump of assembler code for function main:
0x080484b3 <+0>: push %ebp
0x080484b4 <+1>: mov %esp,%ebp
0x080484b5 <+2>: mov 0xc(%ebp),%eax
0x080484b6 <+3>: add $0x4,%eax
0x080484b7 <+4>: mov (%eax),%eax
0x080484b8 <+5>: push %eax
0x080484bf <+12>: eip call 0x804846b <greeting>
0x080484c4 <+17>: eip add $0x4,%esp
```

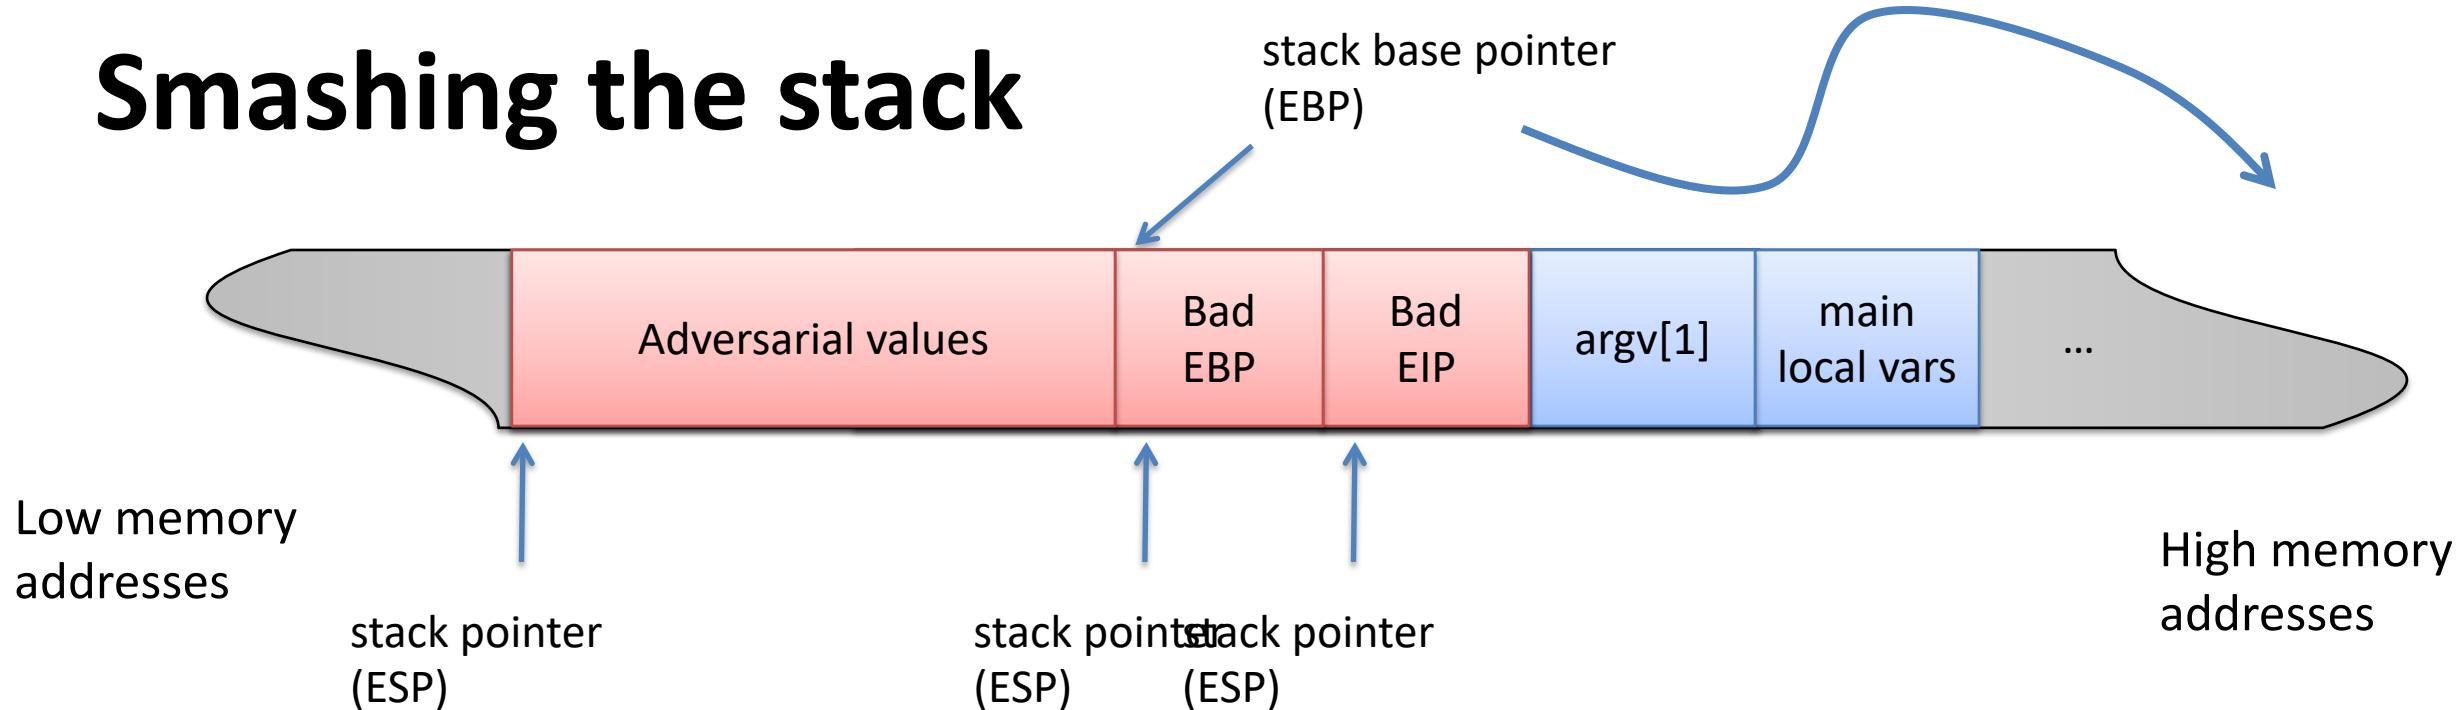
Pushing argv[1] onto stack

```
(gdb) disassemble greeting
Dump of assembler code for function greeting:
0x0804846b <+0> eip push %ebp
0x0804846c <+1> eip mov %esp,%ebp
0x0804846d <+2> eip sub $0x190,%esp
```

.... (more stuff including strcpy) ...

```
0x080484b1 <+70> eip leave
0x080484b2 <+71> eip ret
```

# Smashing the stack



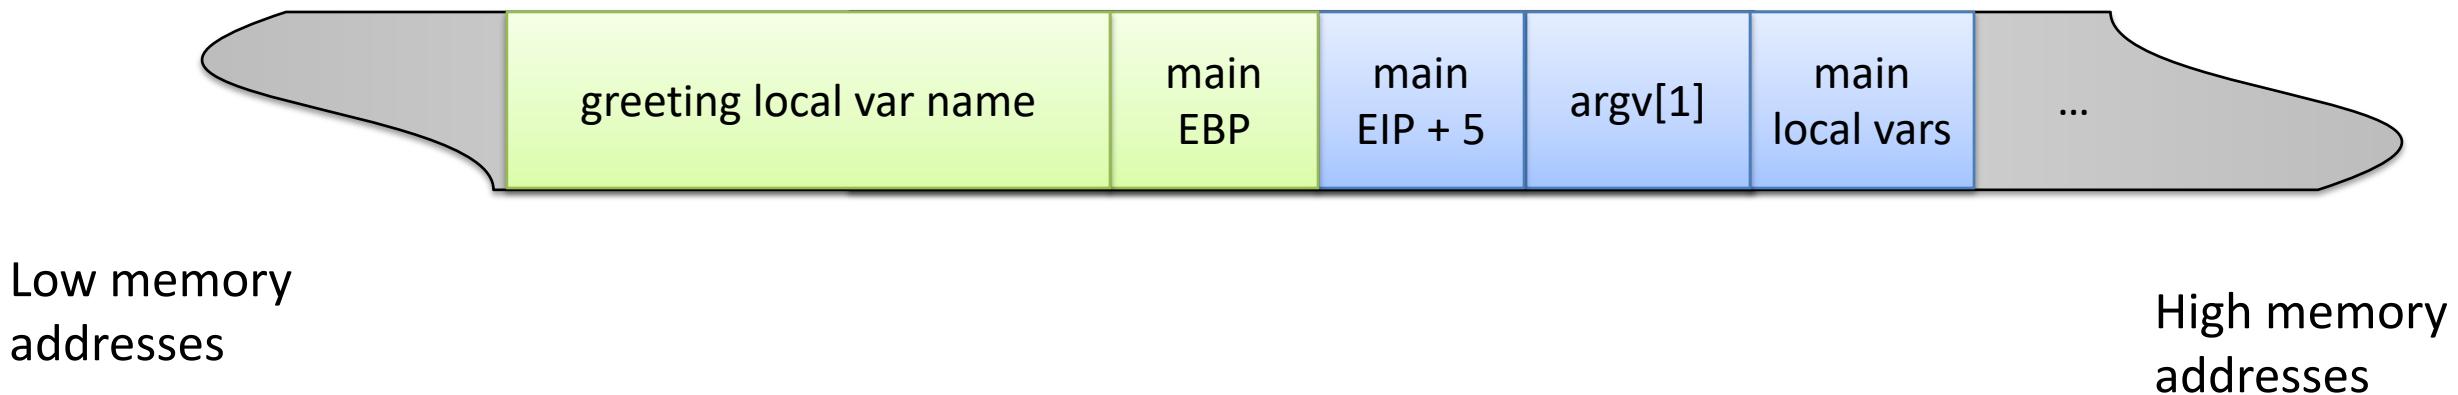
```
student@5435-hw4-vm:~/demo$ gdb -q meet
Reading symbols from meet...done.
(gdb) disassemble main
Dump of assembler code for function main:
0x080484b3 <+0>: push %ebp
0x080484b4 <+1>: mov %esp,%ebp
0x080484b5 <+2>: mov 0xc(%ebp),%eax
0x080484b6 <+3>: add $0x4,%eax
0x080484b7 <+4>: mov (%eax),%eax
0x080484b8 <+5>: push %eax
0x080484b9 <+6>: call 0x804846b <greeting>
0x080484bf <+12>: add $0x4,%esp
0x080484c4 <+17>:
```

Pushing argv[1] onto stack

Bad eip

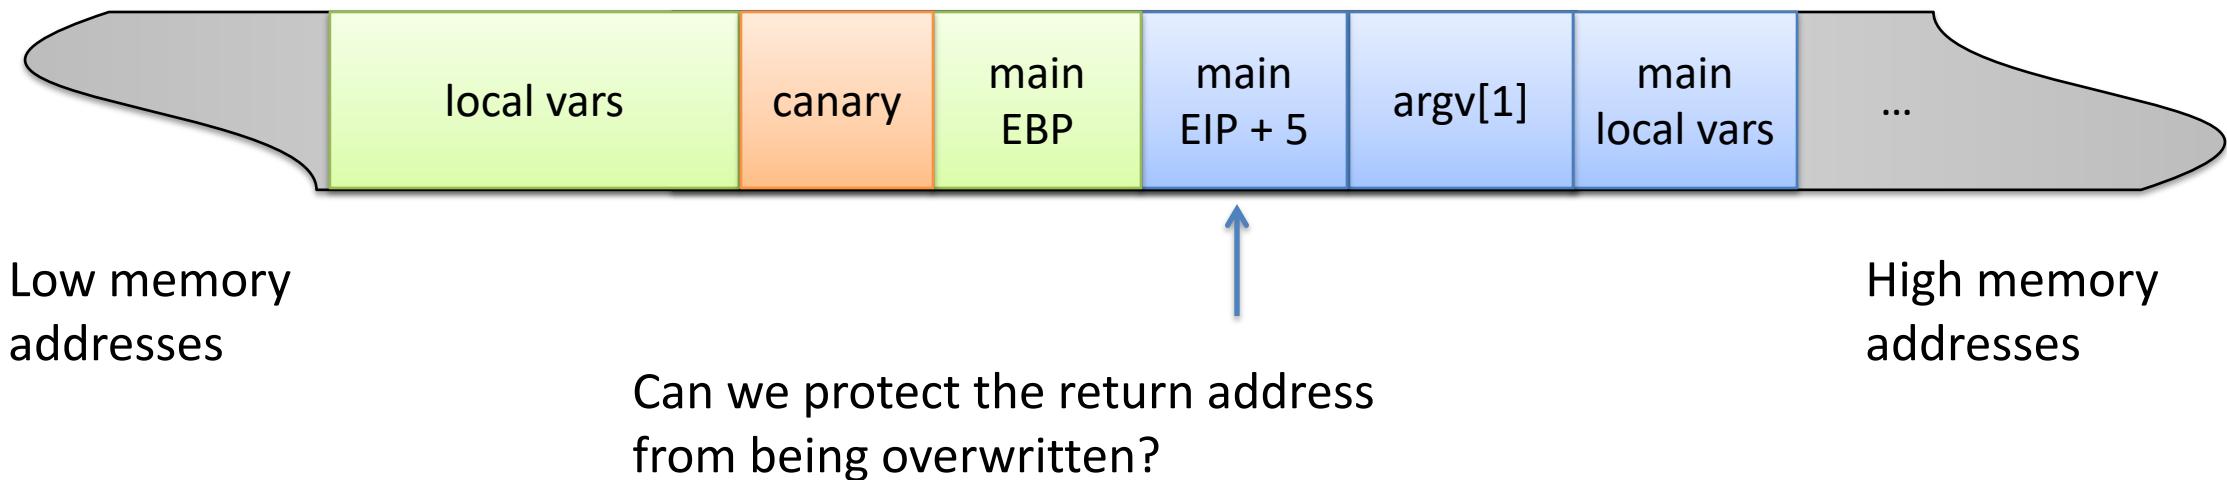
```
(gdb) disassemble greeting
Dump of assembler code for function greeting:
0x0804846b <+0>: push %ebp
0x0804846c <+1>: mov %esp,%ebp
0x0804846e <+3>: sub $0x190,%esp
...
.... (more stuff including strcpy) ...
0x080484b1 <+70> eip leave
0x080484b2 <+71> eip ret
```

# Countermeasures?



- Stack protections
- Address space layout randomization
- W^X (non-executable stack)
- Control-flow integrity

# Detecting buffer overflows: stack canaries



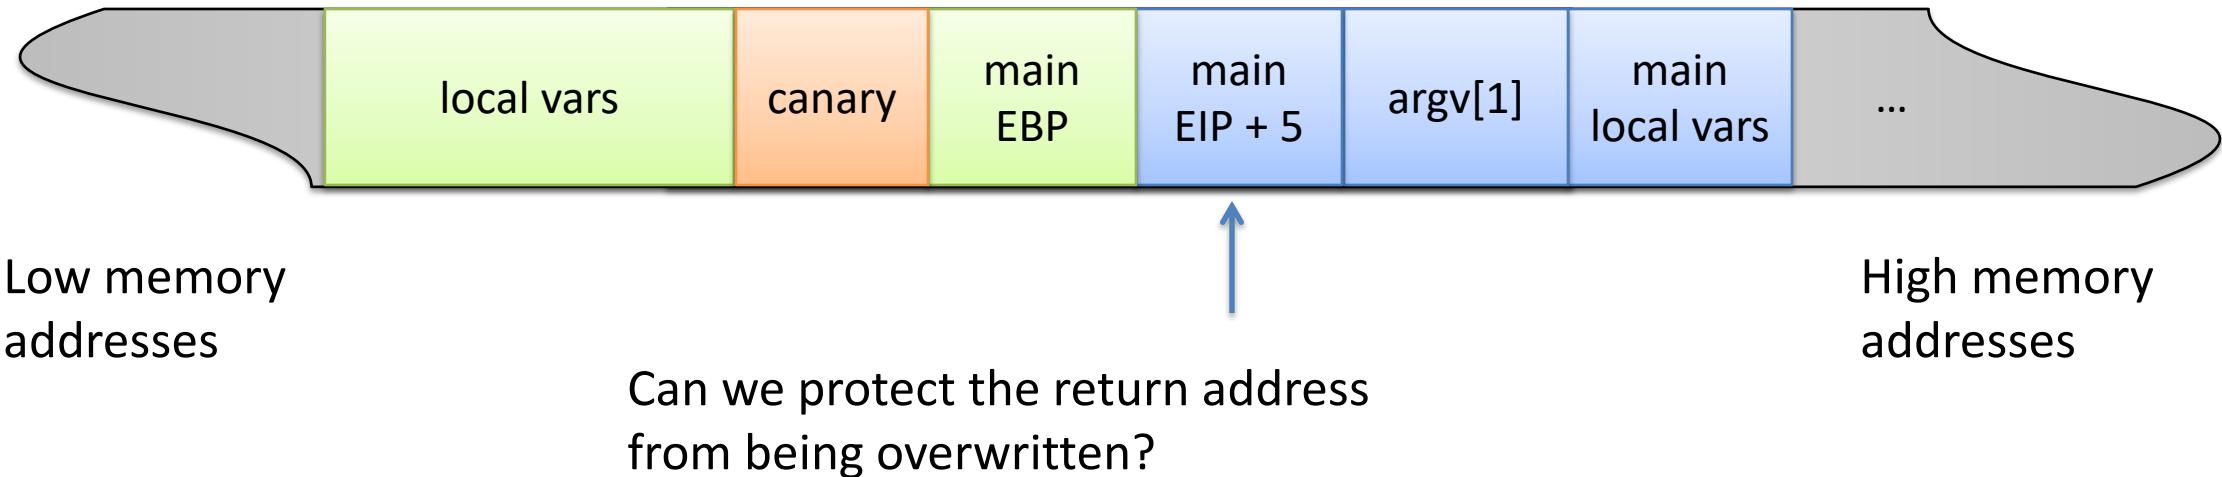
**Before executing function, write canary value to stack**

## Three type:

- Random value (choose once for whole process)
  - Terminator canary (NULL bytes / EOF / etc. so string functions won't copy past canary)
  - Random XOR canary (XOR random value and stored EIP and/or EBP)

On end of function, check that canary is correct, if not terminate process

# Detecting buffer overflows: stack canaries



## StackGuard paper (1998):

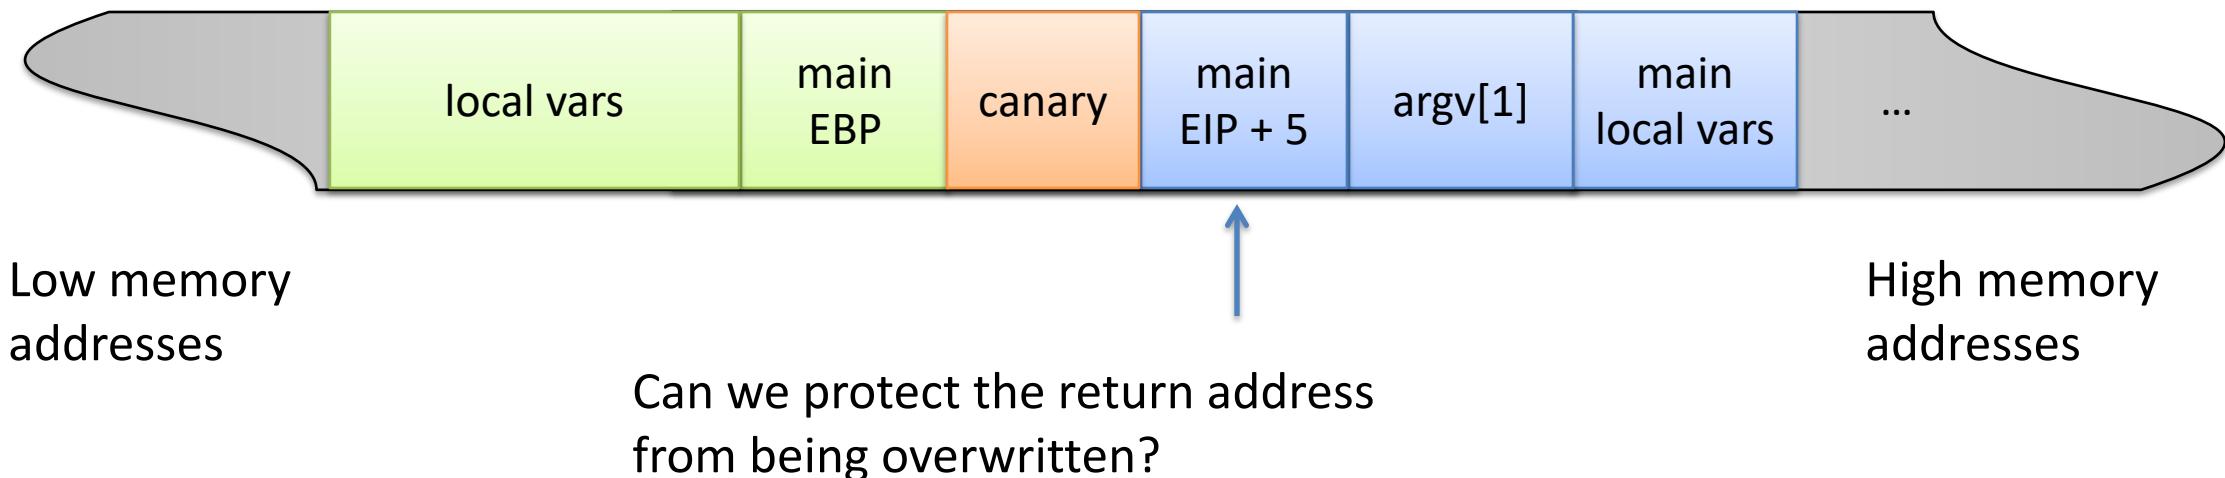
- GCC extension that adds runtime canary checking between saved EIP and EBP
  - 8% overhead on Apache

## ProPolice (IBM gcc patches 2001-2005):

- Moved canary to after saved EBP
  - Rearranges local variables so arrays are highest in stack

## What is the problem with this?

# Detecting buffer overflows: stack canaries



## StackGuard paper (1998):

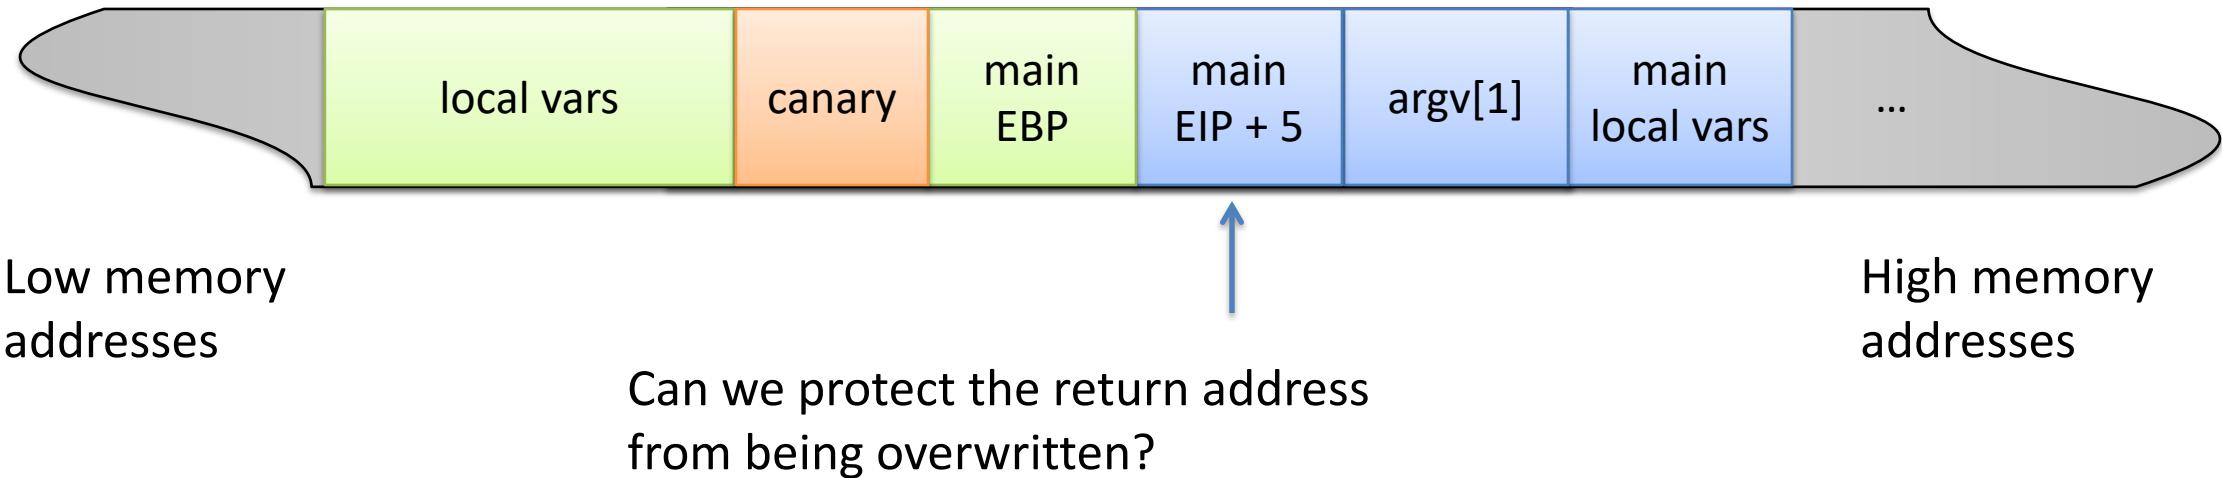
- GCC extension that adds runtime canary checking between saved EIP and EBP
  - 8% overhead on Apache

## ProPolice (IBM gcc patches 2001-2005):

- Moved canary to after saved EBP
  - Rearranges local variables so arrays are highest in stack

# What is the problem with this?

# Detecting buffer overflows: stack canaries



## StackGuard paper (1998):

- GCC extension that adds runtime canary checking between saved EIP and EBP
  - 8% overhead on Apache

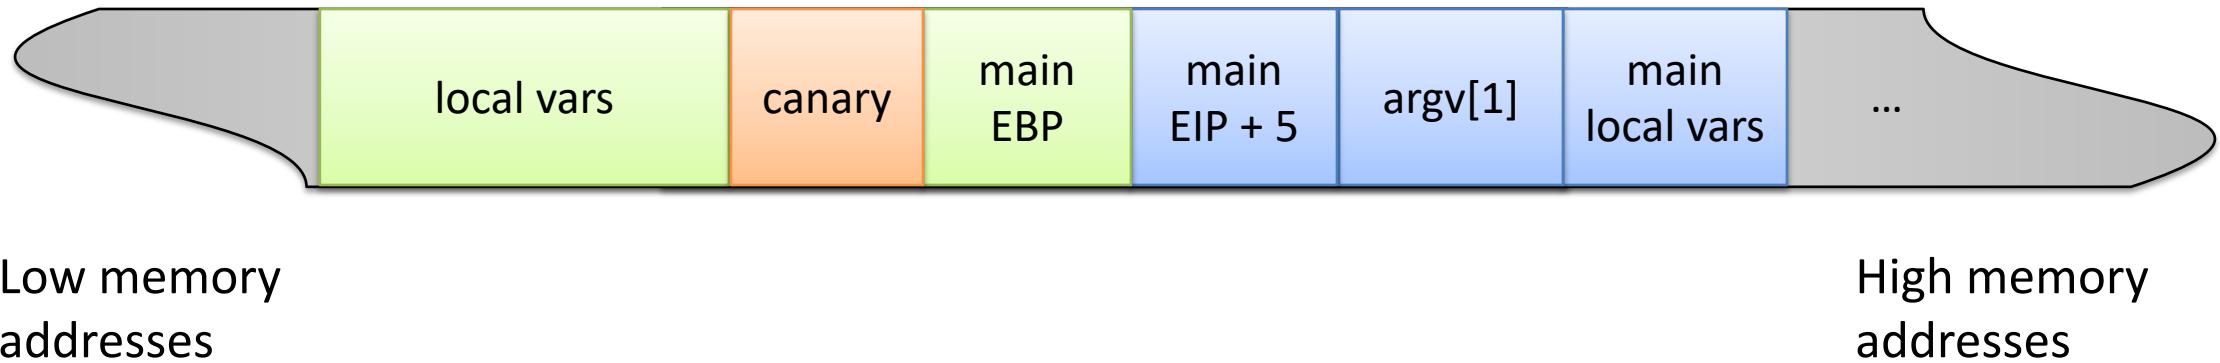
## ProPolice (IBM gcc patches 2001-2005):

# What is the problem with this?

- Moved canary to after saved EBP
- Rearranges local variables so arrays are highest in stack

## Protect local var pointers

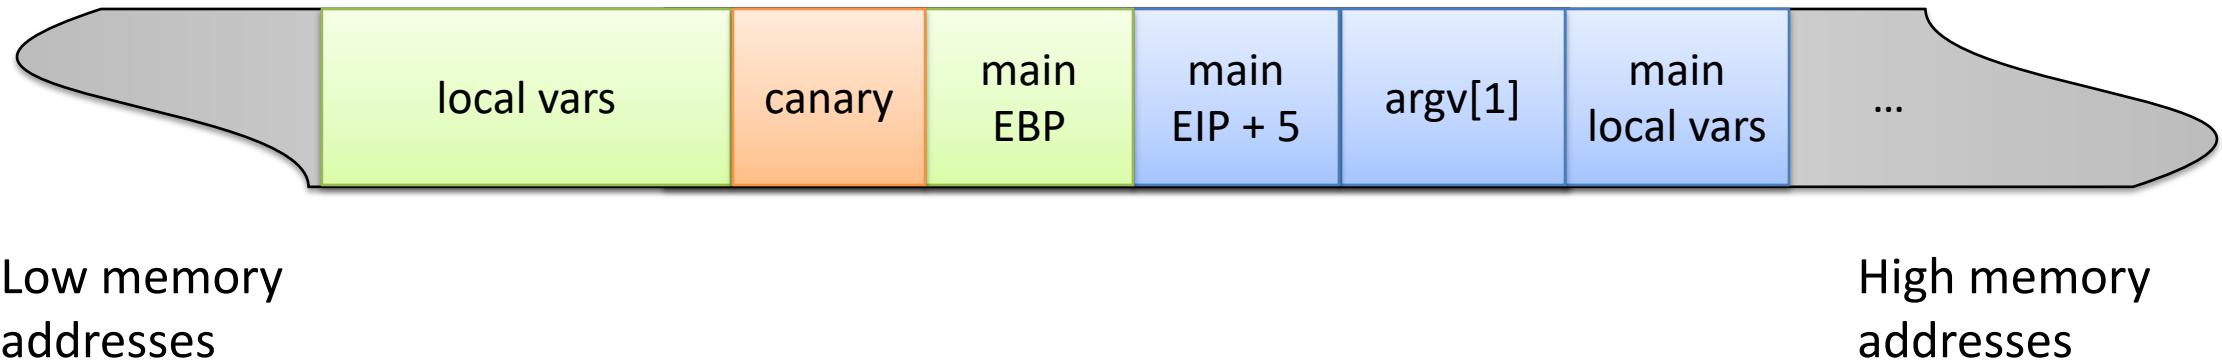
# Detecting buffer overflows: stack canaries



Contemporary gcc implementation

Flag	Default?	Notes
-fno-stack-protector	No	Turns off protections
-fstack-protector	Yes	Adds to funcs that call <code>alloca()</code> & w/ arrays larger than 8 chars ( <code>--param=ssp-buffer-size changes 8</code> )
-fstack-protector-strong	No	Also funcs w/ any arrays & refs to local frame addresses
-fstack-protector-all	No	All funcs

# Detecting buffer overflows: stack canaries



How to circumvent canary protection?

<http://www.phrack.org/issues.html?issue=56&id=5>

```
1 #include <stdio.h>
2 #include <string.h>
3
4
5 int f (char ** argv)
6 {
7     int pipa; // useless variable
8     char *p;
9     char a[30];
10
11    p=a;
12
13    printf ("p=%x\t -- before 1st strcpy\n",p);
14    strcpy(p,argv[1]);           // <== vulnerable strcpy()
15    printf ("p=%x\t -- after 1st strcpy\n",p);
16    strncpy(p,argv[2],16);
17    printf("After second strcpy ;)\n");
18 }
19
20
21 int main(int argc, char* argv[] )
22 {
23     greeting( argv[1] );
24     printf( "Bye %s\n", argv[1] );
25 }
```

# “Reading” the stack to recover canary



Request (can trigger buffer overflow in stack)

Apache forks  
off child process  
to handle request

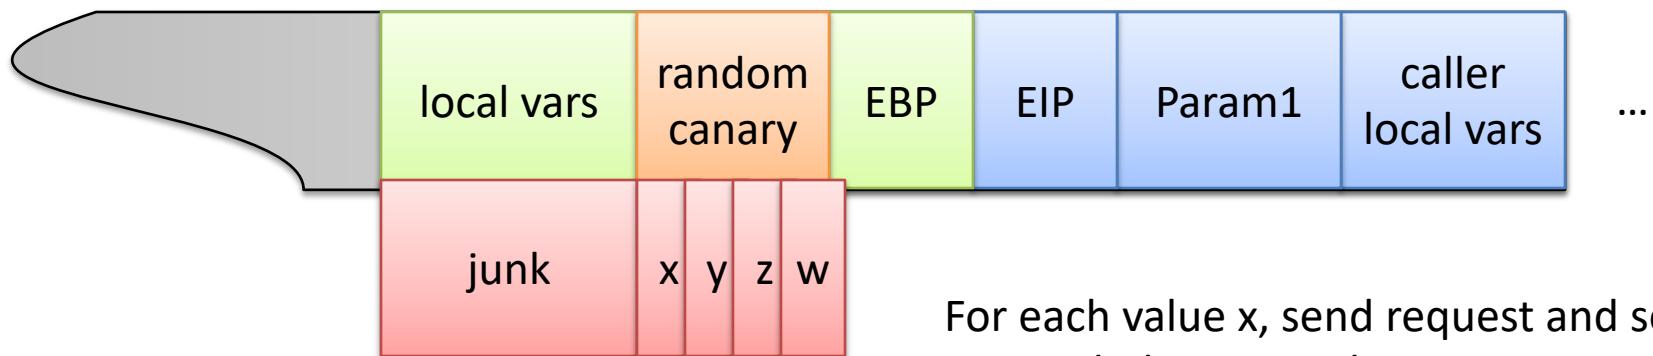


fork() copies process memory into  
child. So every child has:

- same canaries
- same address randomization

Response or process crashes

Apache web server w/ buffer overflow

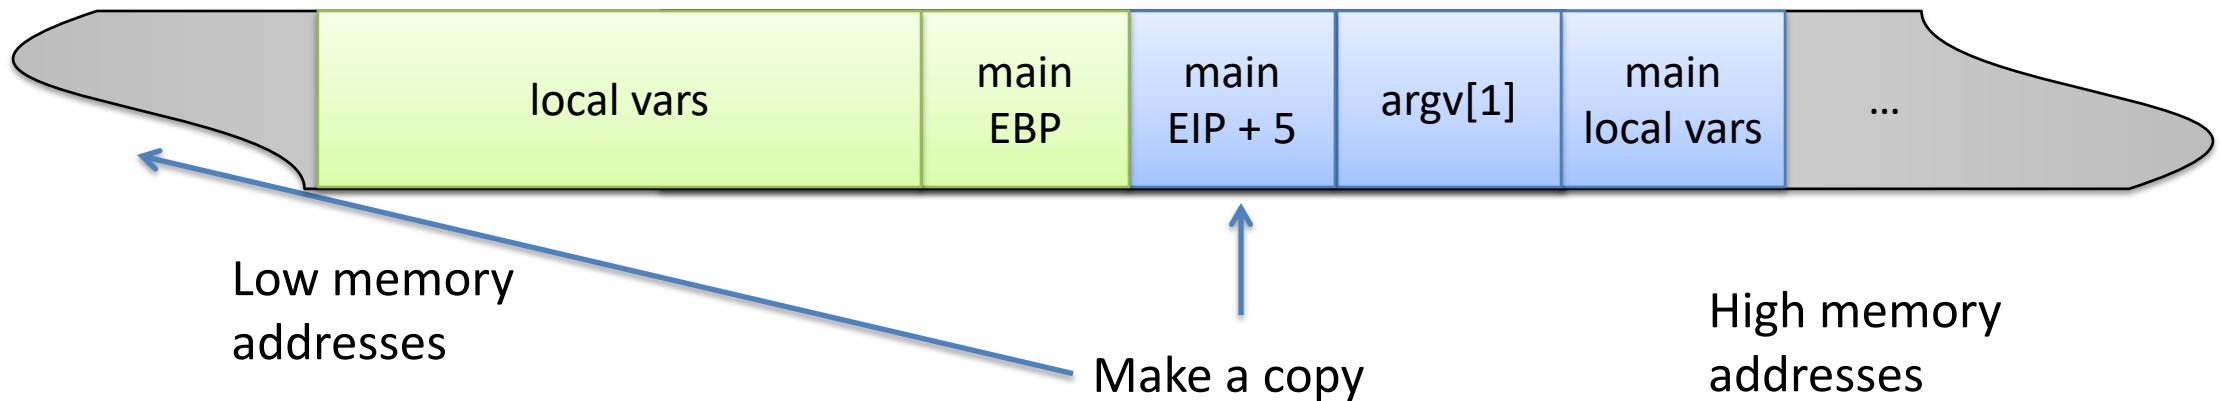


For each value x, send request and see if  
responded to properly

Expected  $2^7 + 2^7 + 2^7 + 2^7 = 512$  requests

Repeat for subsequent bytes of canary

# Detection: copying values to safer location



## StackShield:

- Function call: copy return address to safer location (beginning of .data)
- Check if stack value is different on function exit

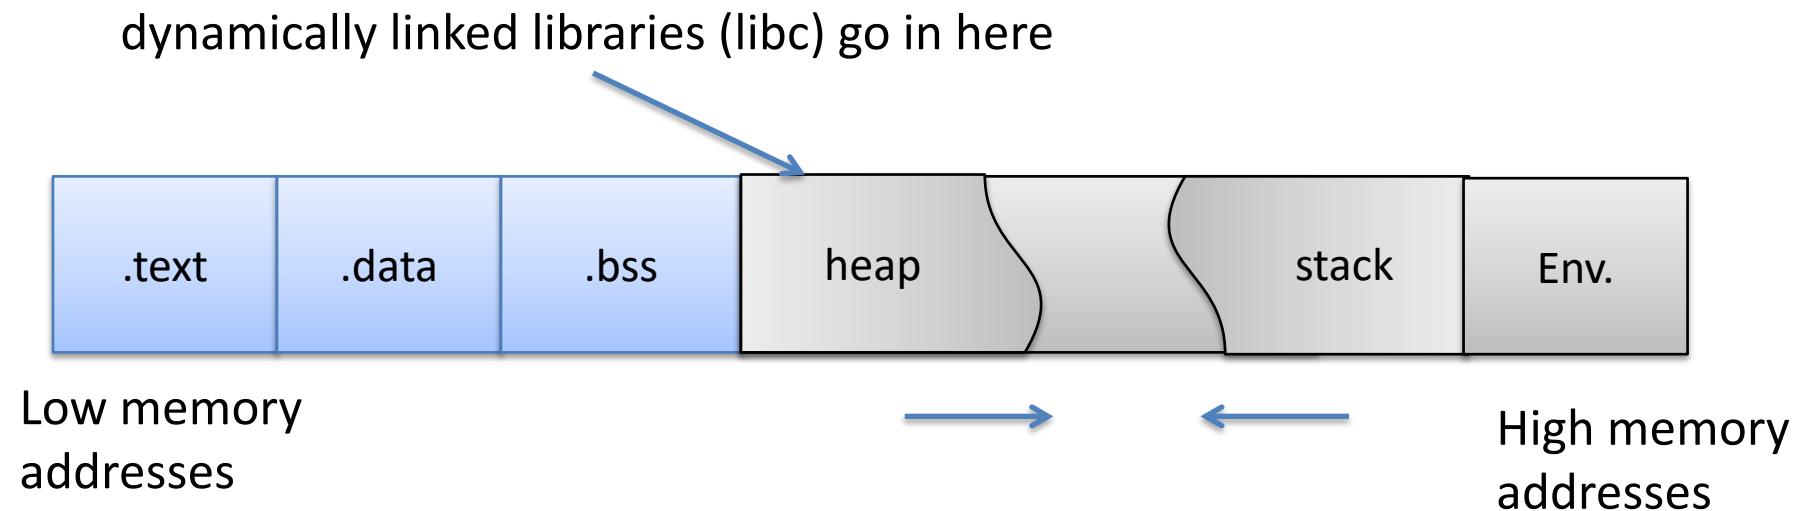
## Circumvention:

- Write over saved copy & return address, if possible
- Hijack control flow without return address

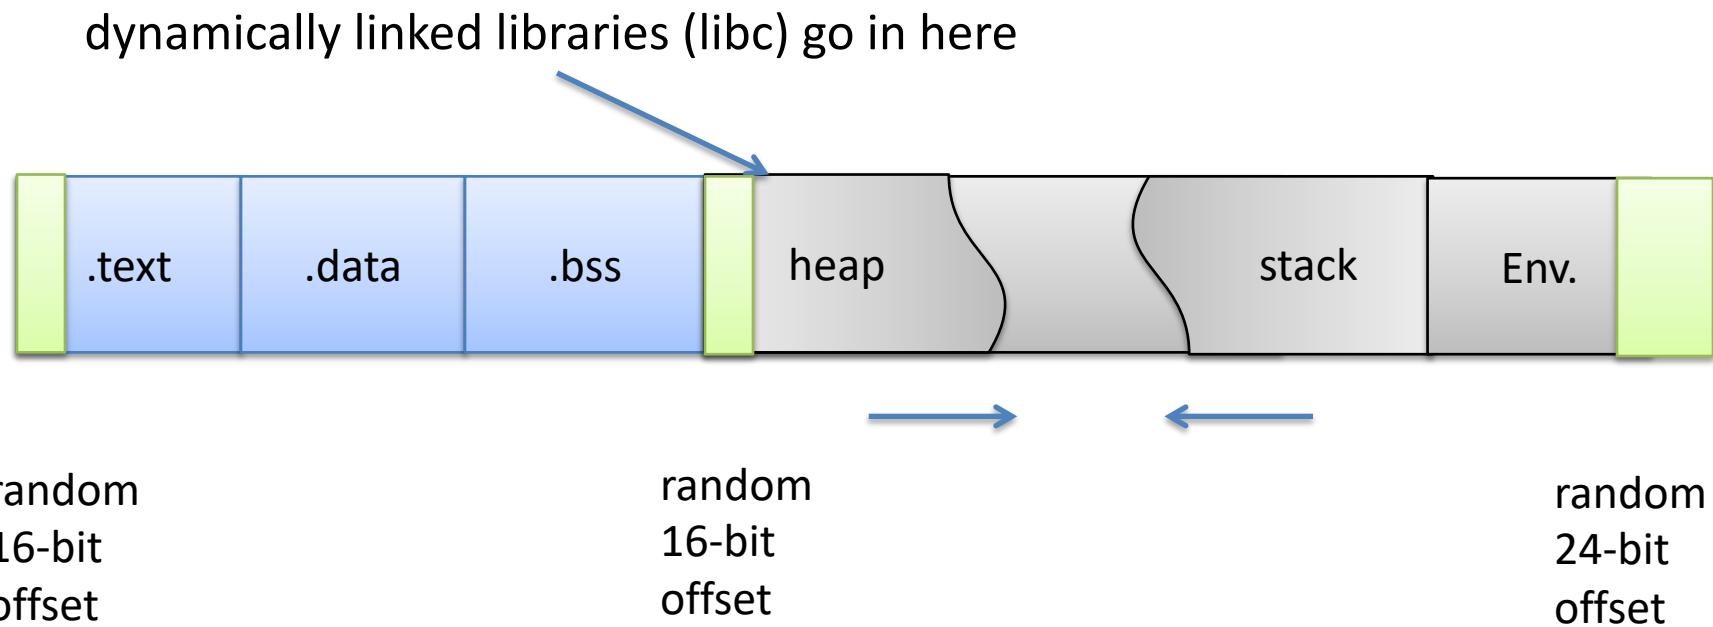
# Countermeasures

- Stack protections
- Address space layout randomization
- W^X (non-executable stack)
- Control-flow integrity

# Address space layout randomization (ASLR)



# Address space layout randomization (ASLR)



Linux PaX implementation for example:

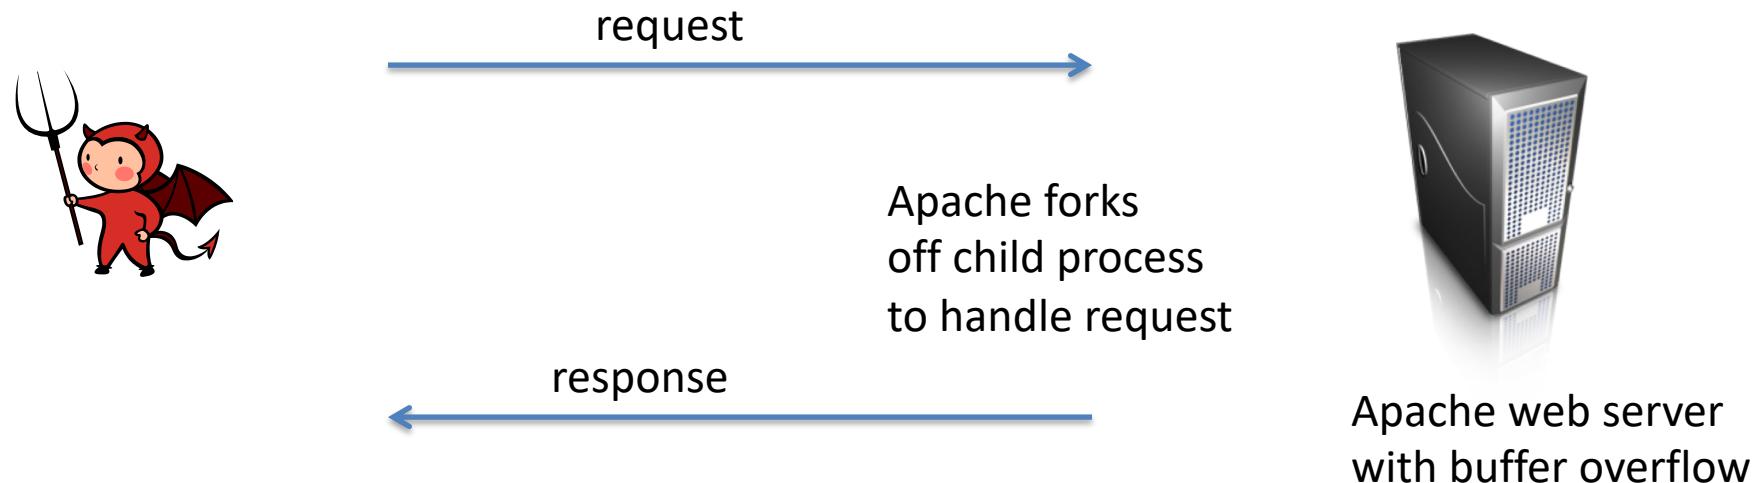
- Randomize offsets of three areas
- 16 bits, 16 bits, 24 bits of randomness. More in 64-bit architectures
- Adds some unpredictability that should help prevent exploits

# Defeating ASLR

- Large nop sled with classic buffer overflow
  - W^X prevents this, stay tuned
- Use a vulnerability that can be used to leak address information
  - E.g., printf arbitrary read
- Brute force the address using forking process

# Defeating ASLR

Brute-forcing example from reading “On the effectiveness of Address Space Layout Randomization” by Shacham et al.



# Defeating ASLR

Brute-forcing example from reading “On the effectiveness of Address Space Layout Randomization” by Shacham et al.



Attacker makes a  
guess of where usleep()  
is located in memory

request

top of stack (higher addresses)
:
0x01010101
0xDEADBEEF
guessed address of usleep()
0xDEADBEEF
64 byte buffer, now filled with A's
:
bottom of stack (lower addresses)



Apache web server  
with buffer overflow

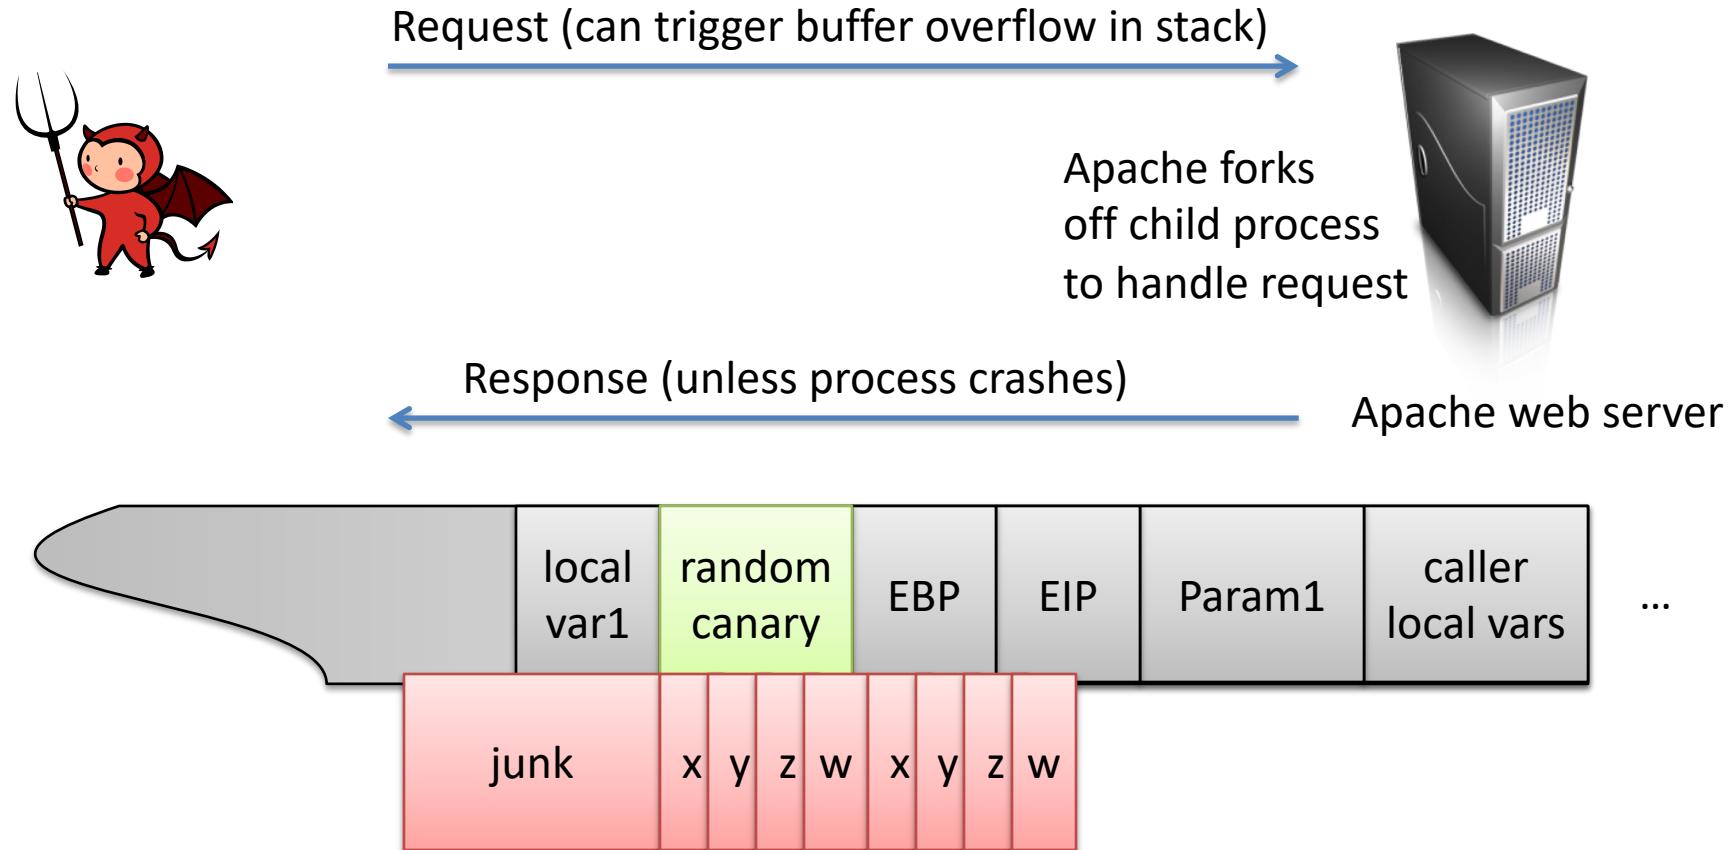
Figure 2: Stack after one probe

Failure will crash the child process  
immediately and therefore kill connection

Success will crash the child process  
after sleeping for 0x01010101  
Microseconds and kill connection

If on 64-bit architecture, such brute-force attack unlikely to work

# Reading the stack, remotely



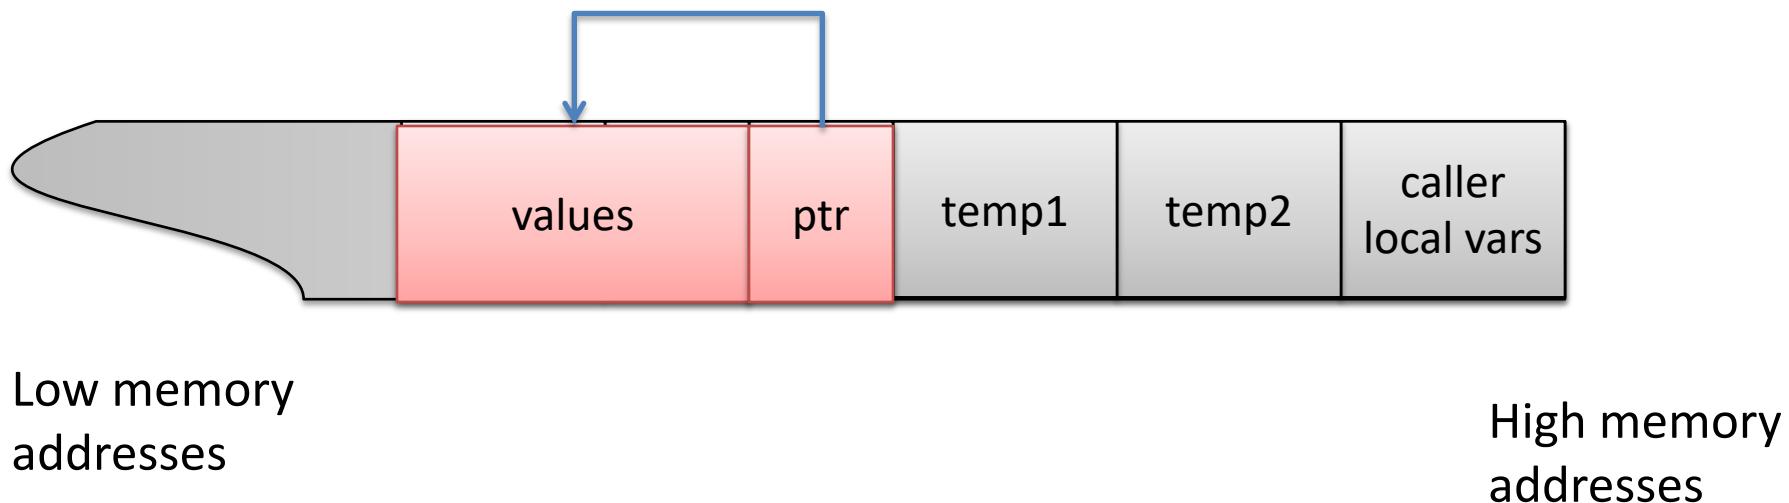
Reading stack for EBP/EIP can give approximate address offset

# Countermeasures

- Stack protections
- Address space layout randomization
- W^X (non-executable stack)
- Control-flow integrity

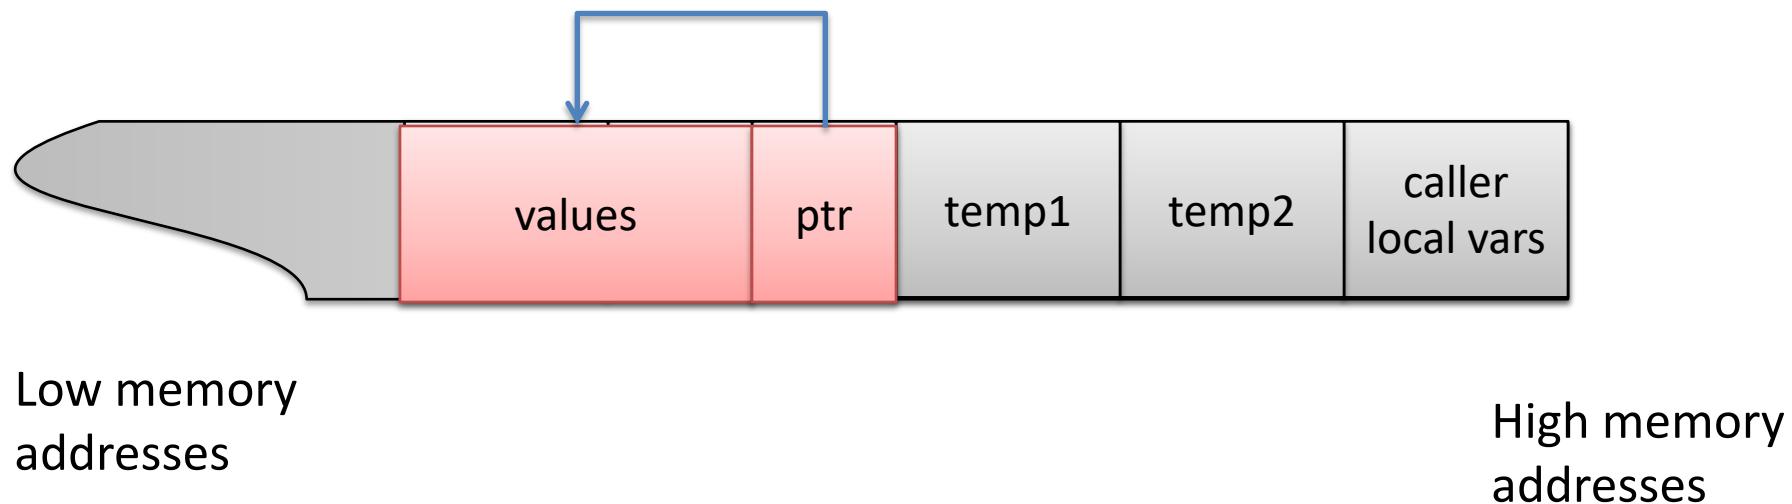
# W^X (W xor X)

- The idea: mark memory page as either
  - Writable or Executable (not both)
- Specifically: make heap and stack non-executable
- Default in gcc now, turn off with `-z execstack`



# **W^X (W xor X)**

- AMD64: NX bit (Non-Executable)  
IA-64: XD bit (eXecute Disabled)  
ARMv6: XN bit (eXecute Never)
  - Extra bit in each page table entry
  - Processor refuses to execute code if bit = 1
  - Mark heap and stack segments as such



# W^X (W xor X)

Software emulation of NX bits

- ExecShield (RedHat Linux)
- PaX (Page-eXec) (uses NX bit if available)

mprotect()

- Process can set permissions on memory pages

# Will W^X prevent:

- AlephOne's stack overflow exploit? Yes
- Stack smash that overwrites pointer to point at shell code in Heap or Env variable? Yes
- Heap overflow with same shell location? Yes
- Double free with same shell location? Yes

# Return-into-libc exploits

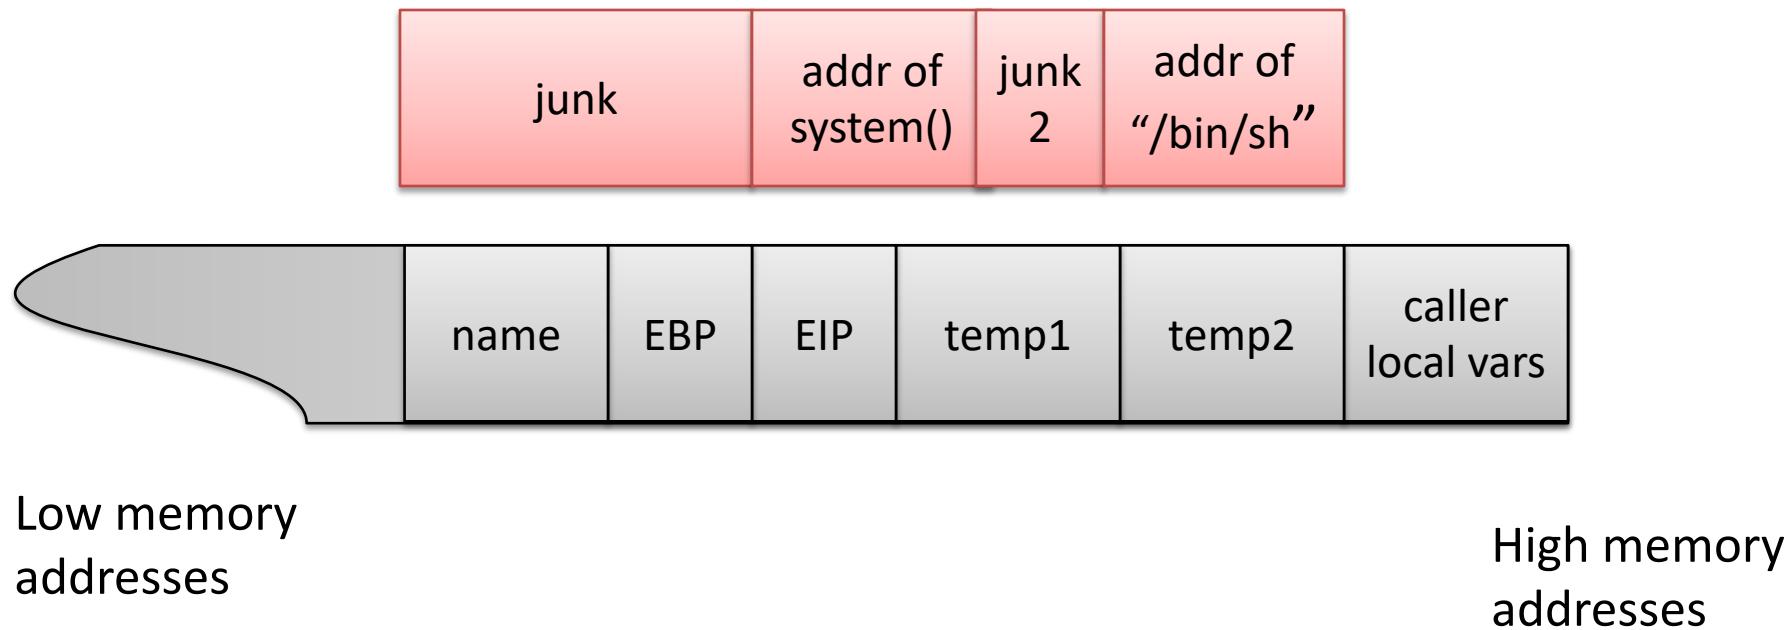
- libc is standard C library, included in all processes
- system() --- execute commands on system

```
(gdb) b main
Breakpoint 1 at 0x80484a0: file sploit1.c, line 15.
(gdb) r
Starting program: /home/user/pp1/sploits/sploit1

Breakpoint 1, main () at sploit1.c:15
15      args[0] = TARGET;
(gdb) p system
$1 = {<text variable, no debug info>} 0xb7ecf180 <system>
(gdb) _
```

# Return-into-libc exploits

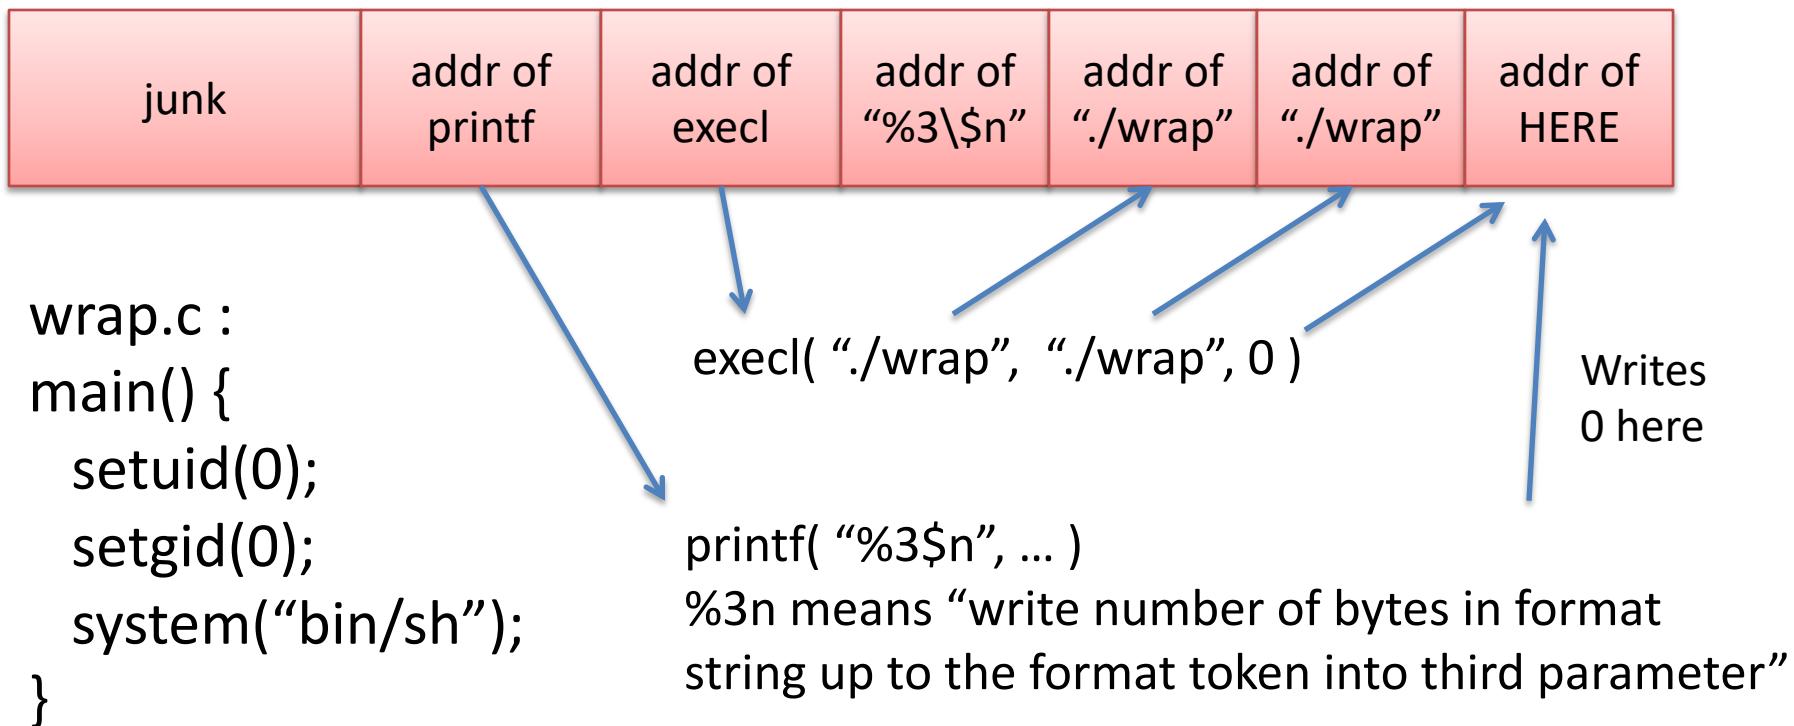
Overwrite EIP with address of system() function  
junk2 just some filler: returned to after system call  
first argument to system() is ptr to “/bin/sh”



# Return-into-libc exploits

This simple exploit has a few deficiencies (for attacker):

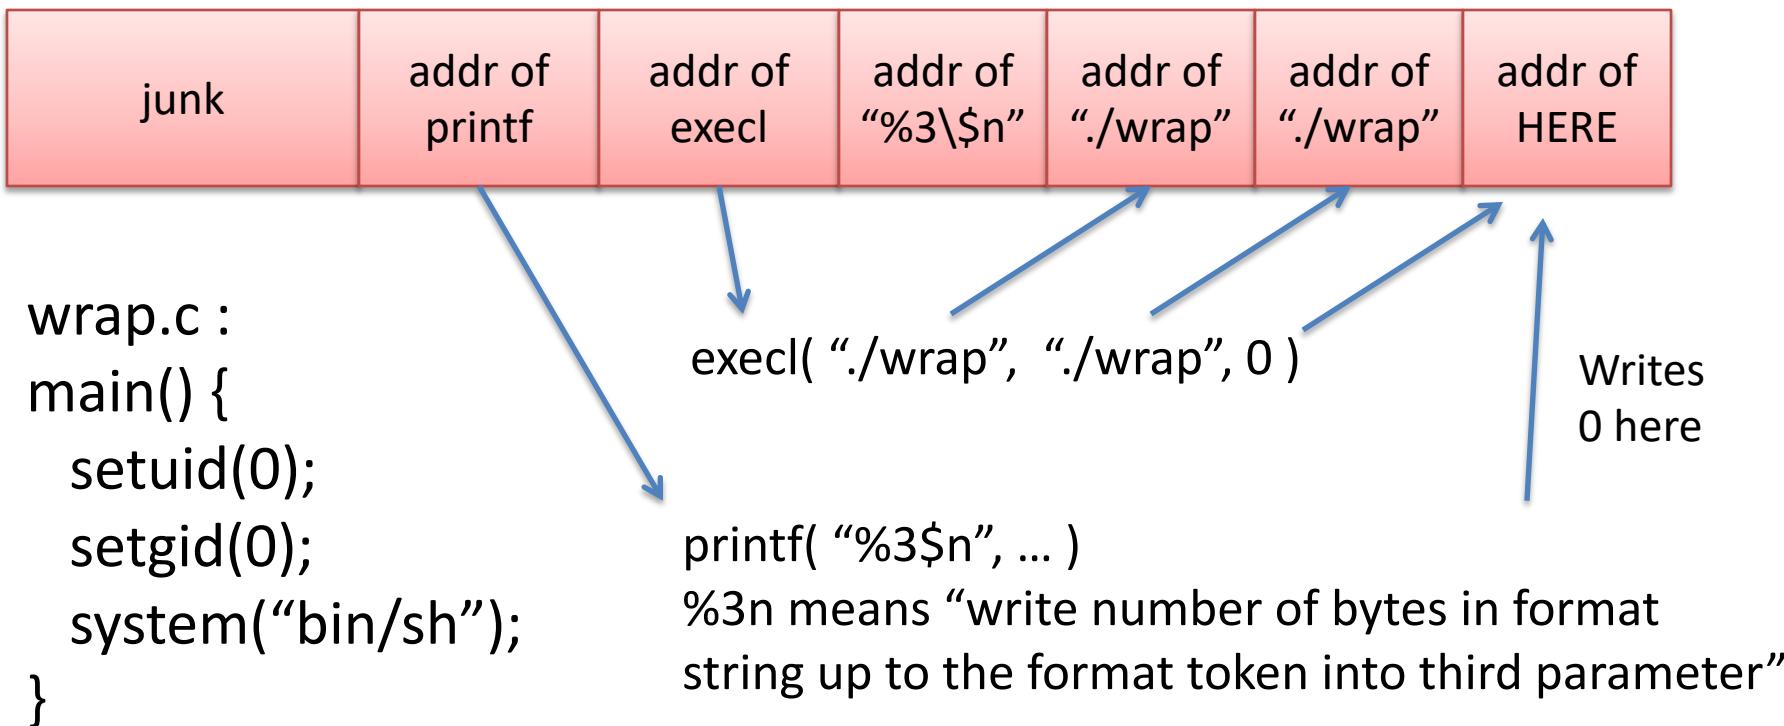
- Crashes after exiting called /bin/sh ( easy to fix with exit() )
- system() drops privileges by default



# Return-into-libc exploits

These exploits only execute instructions marked executable

W^X cannot stop such an attack



# Return-into-libc exploits

Return-into-libc may seem limited:

- Only useful for calling libc functions
- Okay in last example, but not always sufficient
- Before W^X, exploit could run arbitrary code

Can we ***not*** inject any malicious code and yet have an exploit that runs ***arbitrary*** code?

# Return-oriented programming (ROP)

- Second return-into-libc exploit:
  - self-modifying exploit buffer to call a sequence of libc calls
- ROP chains together sequences of calls to “gadgets”
  - Gadget = sequence of instructions that do something useful
- Can find gadgets in executable code pages (e.g., stdlib)
- Shacham showed that stdlib is Turing-complete

# Countermeasures

- Stack protections
- Address space layout randomization
- W^X (non-executable stack)
- Control-flow integrity

# Control-flow integrity

- Abadi et al. 2005
- Recall: control flow graph is (static) set of paths in program
- Include checks that call is found in control-flow graph
  - Simplest: memory locations callable or not (one bit)
  - More complex: label memory locations with ID, use to match-up caller-callee edges in CFG
- Versions implemented in Microsoft Visual Studio, clang

# Countermeasures

- Stack protections
- Address space layout randomization
- W^X (non-executable stack)
- Control-flow integrity

Arms race between exploit designers and countermeasures

- Much harder to exploit some vulnerabilities
- Costs more to develop exploits

Next time: finding vulnerabilities in software

# Return-oriented programming (ROP)

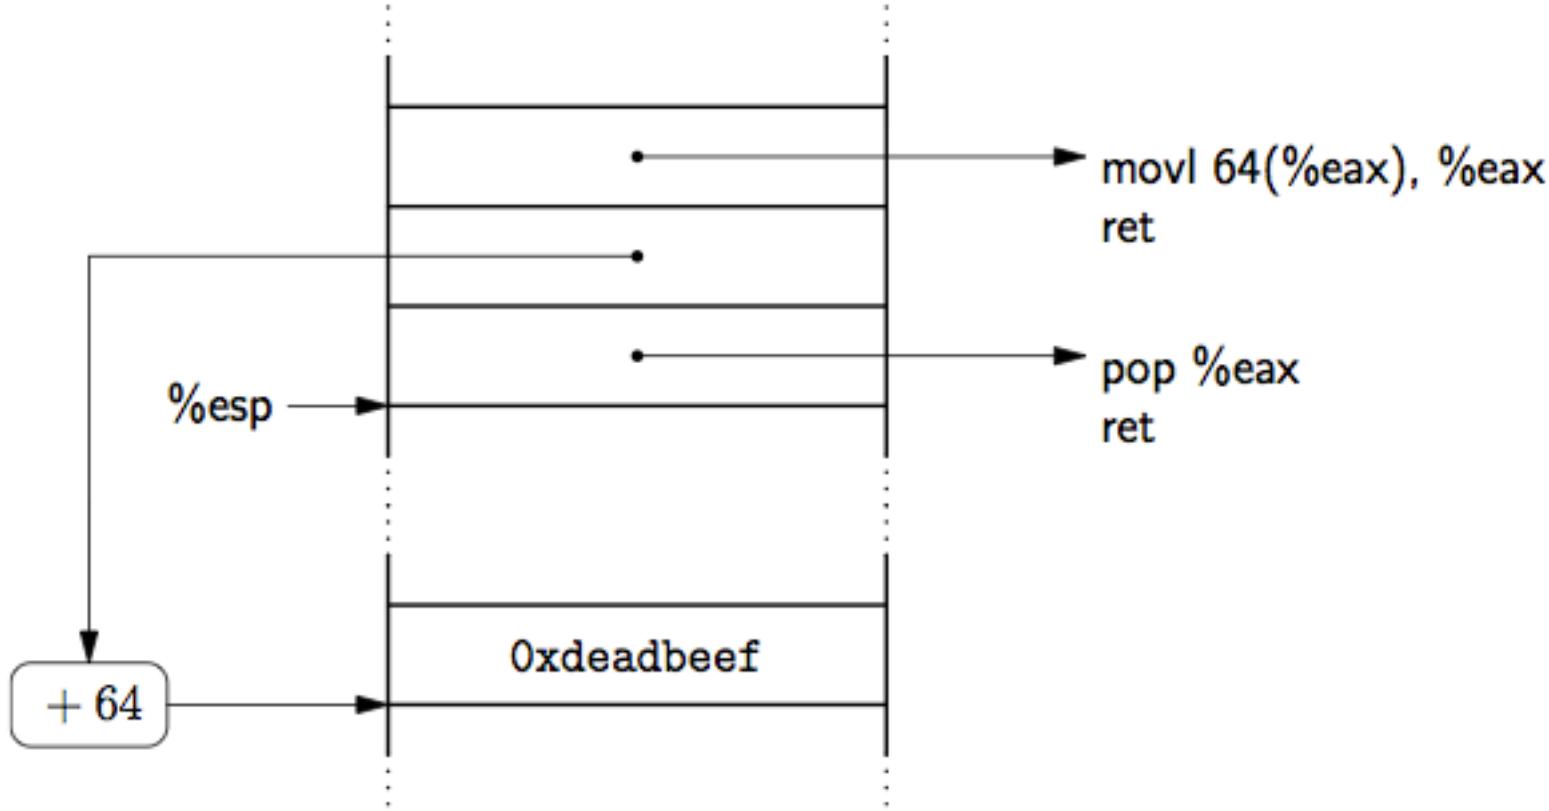
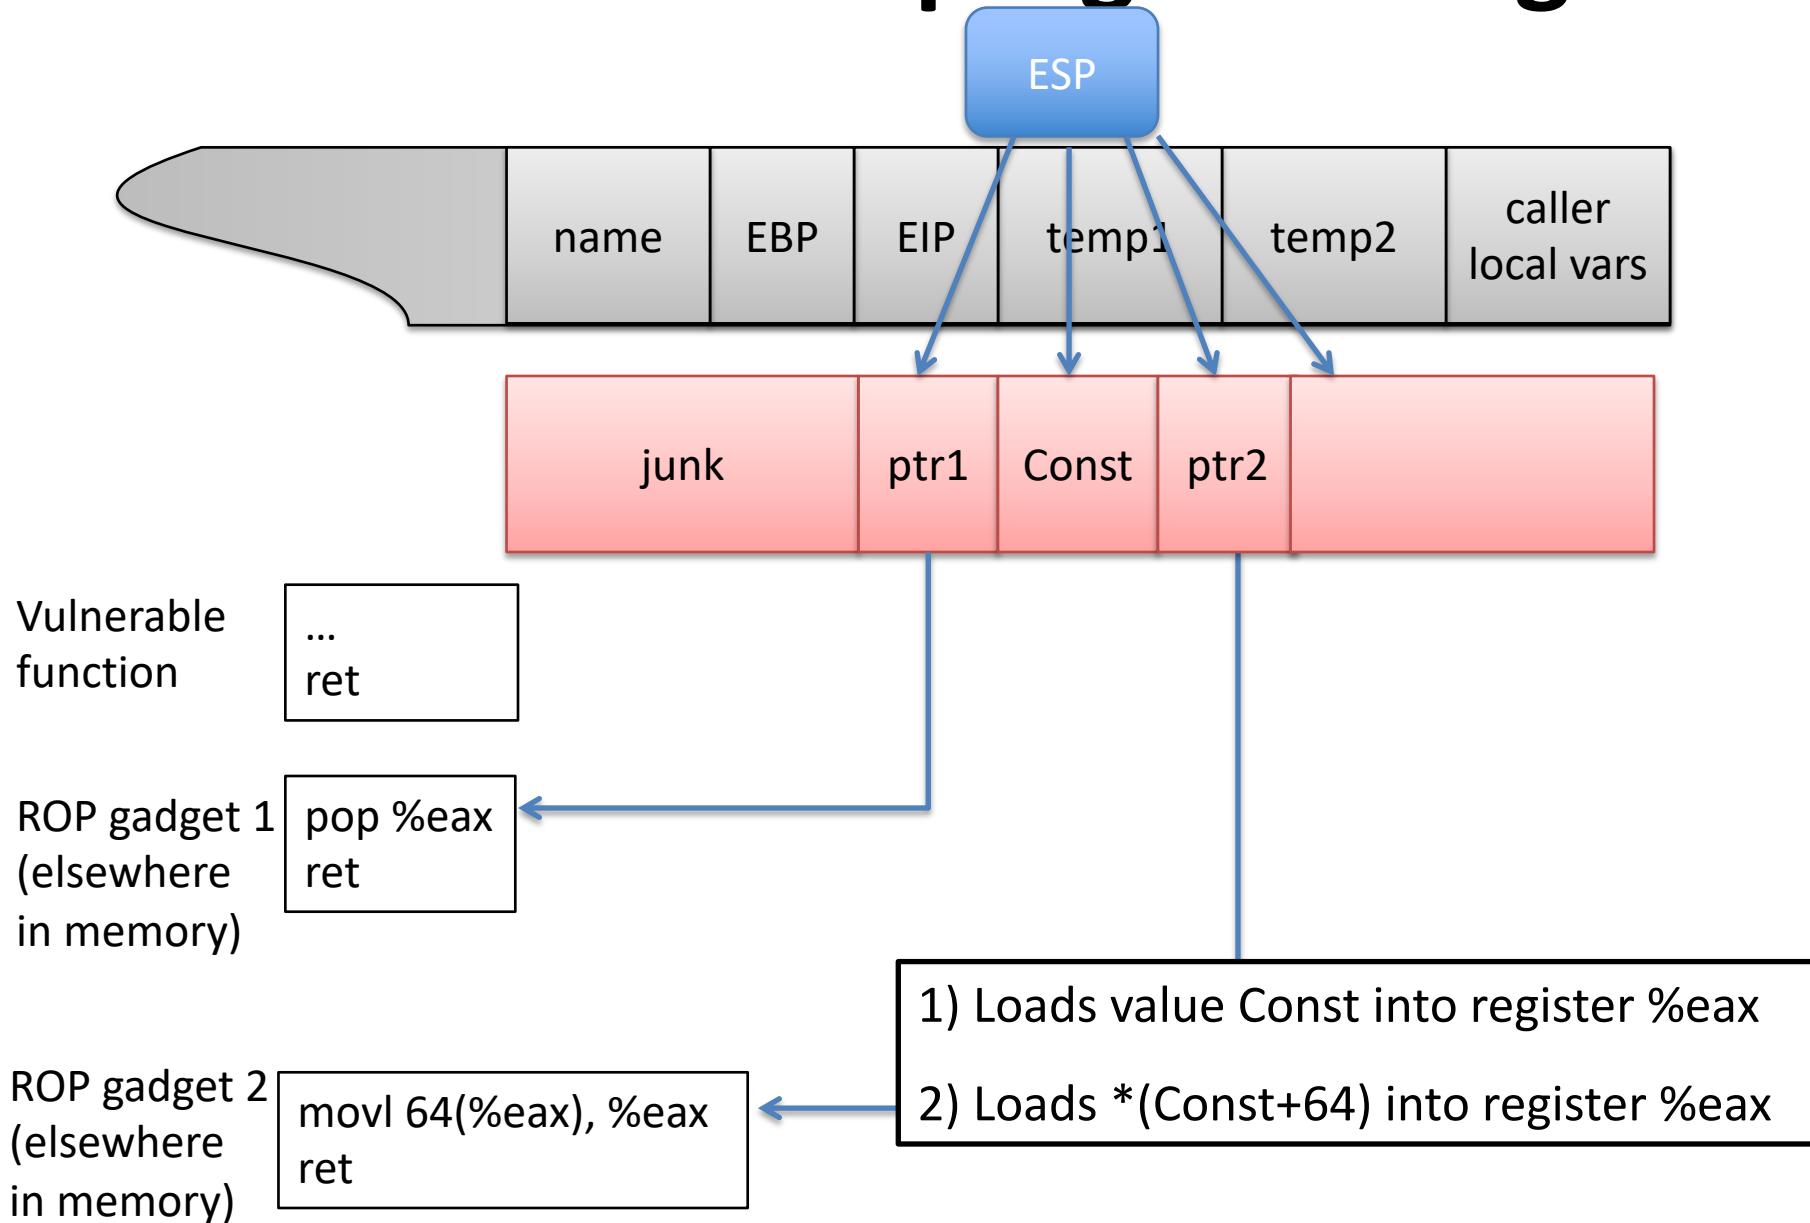


Figure 3: Load a word in memory into `%eax`.

From Shacham “The Geometry of Innocent Flesh on the Bone...” 2007

# Return-oriented programming



# Return-oriented programming (ROP)

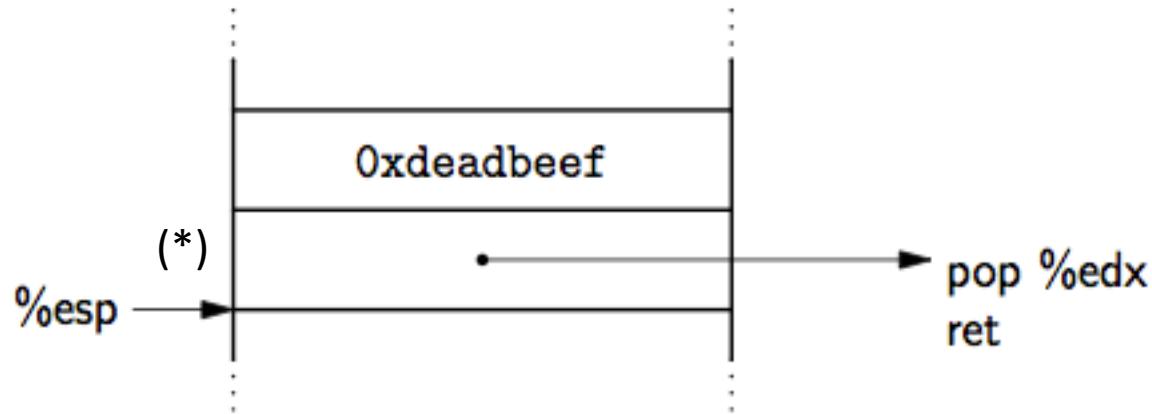


Figure 2: Load the constant **0xdeadbeef** into **%edx**.

From Shacham “The Geometry of Innocent Flesh on the Bone...” 2007

If this is on stack and (\*) is return pointer after buffer overflow, then the result will be loading 0xdeadbeef into edx register

From  
Shacham  
“The Geometry of  
Innocent Flesh on  
the Bone...” 2007

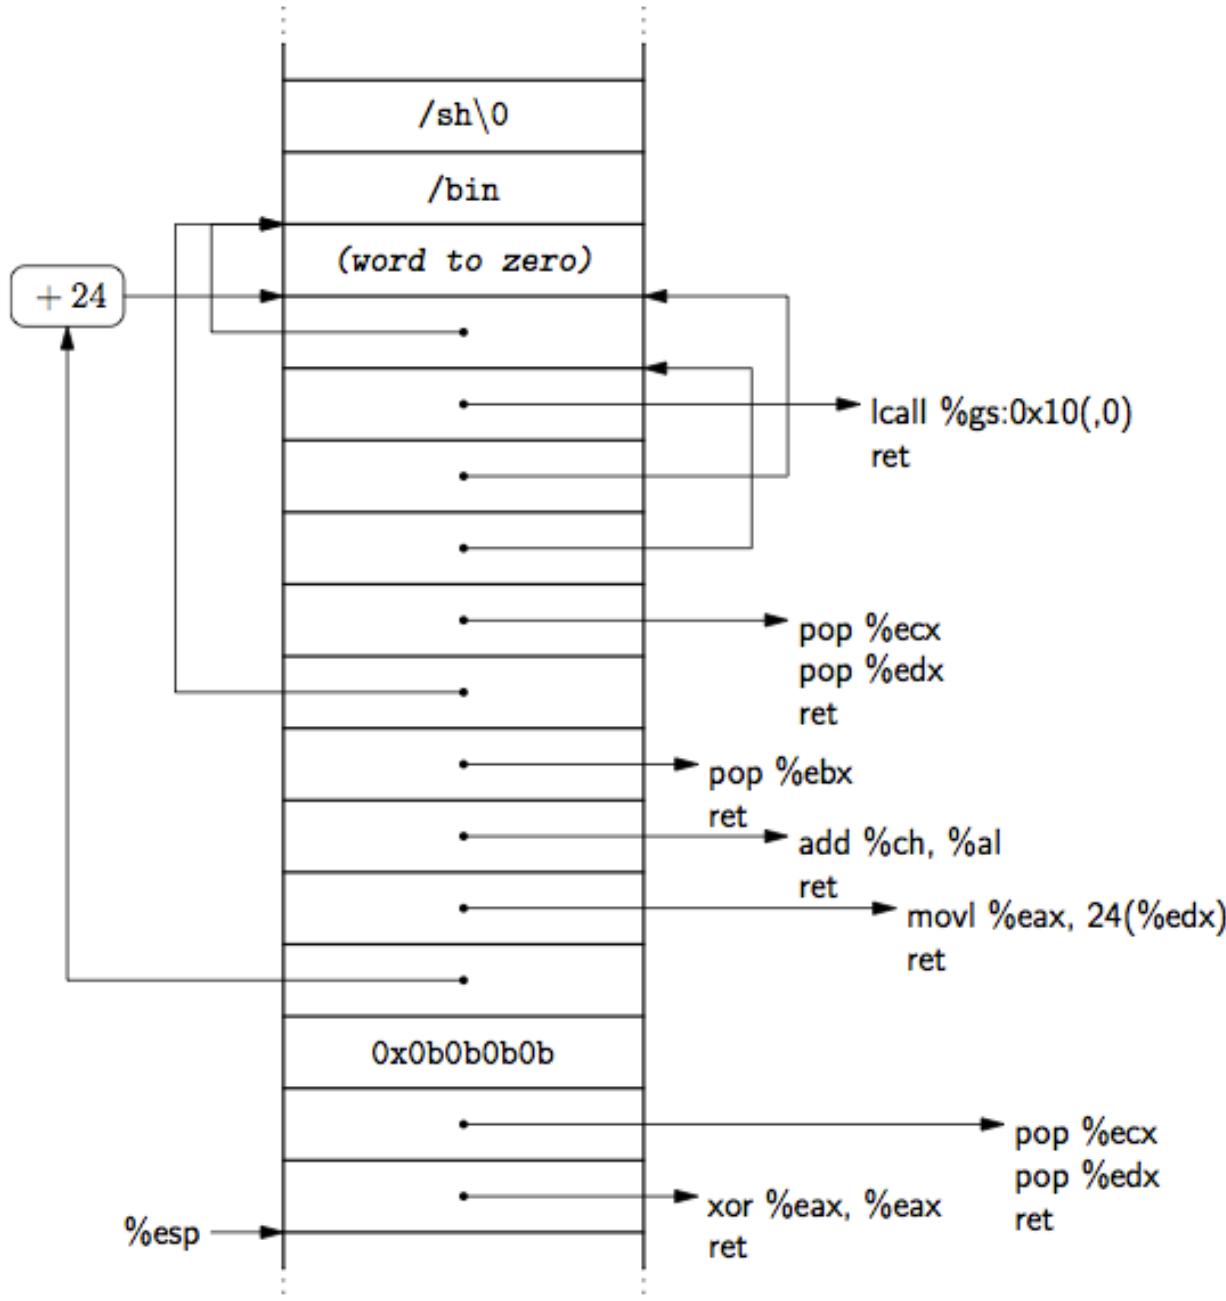


Figure 16: Shellcode.

# ROP where do we get code snippets?

